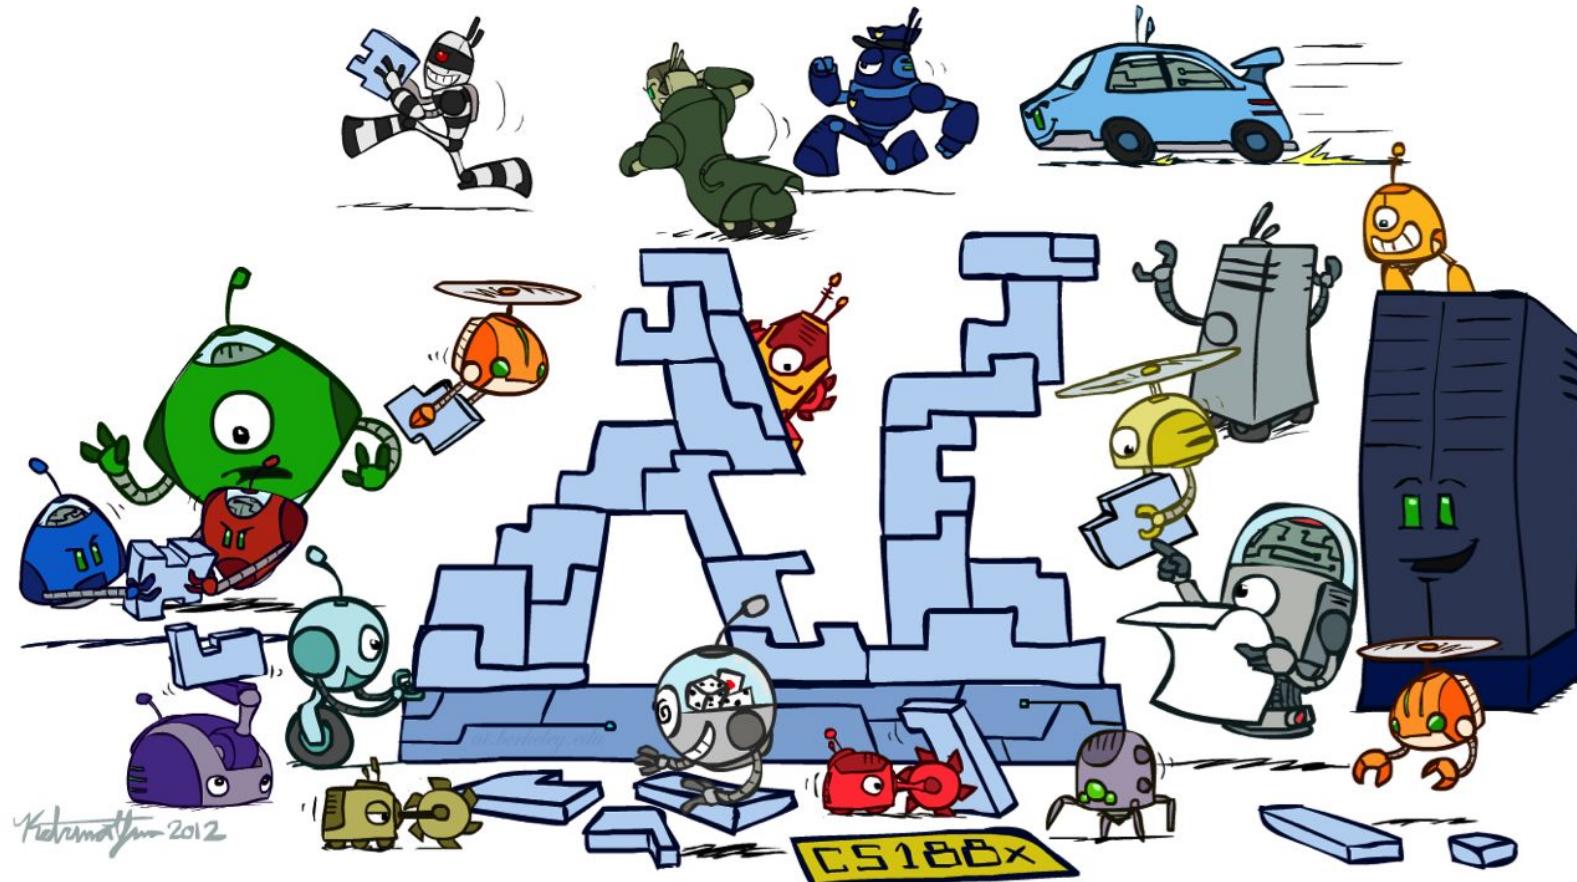


# CS 188: Artificial Intelligence

## Large Language Models



Instructors: Peyrin Kao and Emma Pierson --- University of California, Berkeley

[These slides were created by Cam Allen, Michael Cohen, Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley.

All CS188 materials are available at <http://ai.berkeley.edu>.]

# Next few classes

---

- Today: LLMs
- Tuesday (11/25): Applications: AI for Healthcare (and extra credit!)
- Thursday (11/27): No class (Thanksgiving!)
- Tuesday (12/02): AI for Equality and Course Wrap-Up

# Today's AI

= ChatGPT 4 ▾ 



**How can I help you today?**

 Message ChatGPT... 

ChatGPT can make mistakes. Consider checking important information.

A\ Untitled ▾ 



**What can I help you with today?**

 Message Claude... 

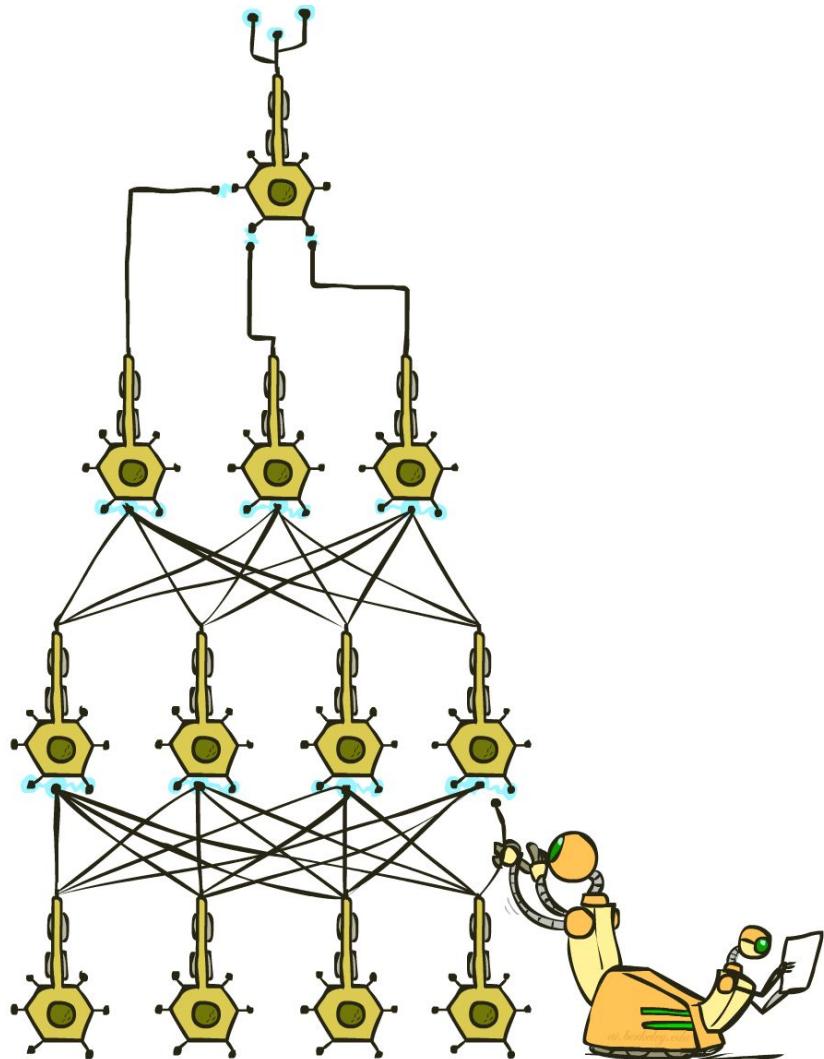
Claude 3 Opus ▾

# Large Language Models

---

- Feature engineering
  - Text tokenization
  - Word embeddings
- Deep neural networks
  - Autoregressive models
  - Self-attention mechanisms
  - Transformer architecture
- Multi-class classification
- Supervised learning
  - Self-supervised learning
  - Instruction tuning
- Reinforcement learning
  - ... from human feedback (RLHF)

# Deep Neural Networks



- Input: some text
  - “The dog chased the”
- Output: more text
  - ... “ ball”
- Implementation:
  - Linear algebra
  - How??

# Text Tokenization

GPT-3.5 & GPT-4    GPT-3 (Legacy)

Many words map to one token, but some don't: `indivisible`.

Unicode characters like emojis may be split into many tokens containing the underlying bytes: `👉`

Sequences of characters commonly found next to each other may be grouped together: `1234567890`

[Clear](#)

[Show example](#)

**Tokens**

**57**

**Characters**

**252**

# Text Tokenization

GPT-3.5 & GPT-4

GPT-3 (Legacy)

Many words map to one token, but some don't: `indivisible`.

Unicode characters like emojis may be split into many tokens containing the underlying bytes: `??????`

Sequences of characters commonly found next to each other may be grouped together: `1234567890`

Text

Token IDs

Tokens

57

Characters

252

# Text Tokenization

GPT-3.5 & GPT-4

GPT-3 (Legacy)

```
[8607, 4339, 2472, 311, 832, 4037, 11, 719, 1063, 1541, 956, 25, 3687,  
23936, 382, 35020, 5885, 1093, 100166, 1253, 387, 6859, 1139, 1690,  
11460, 8649, 279, 16940, 5943, 25, 11410, 97, 248, 9468, 237, 122, 271,  
1542, 45045, 315, 5885, 17037, 1766, 1828, 311, 1855, 1023, 1253, 387,  
41141, 3871, 25, 220, 4513, 10961, 16474, 15]
```

Text

Token IDs

Tokens

57

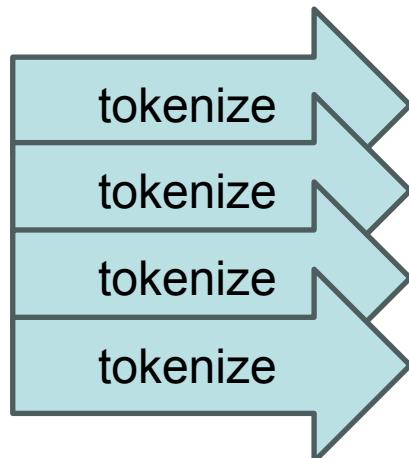
Characters

252

# Word Embeddings

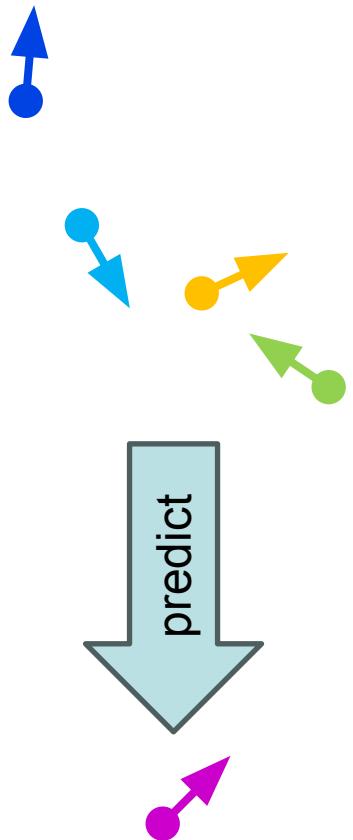
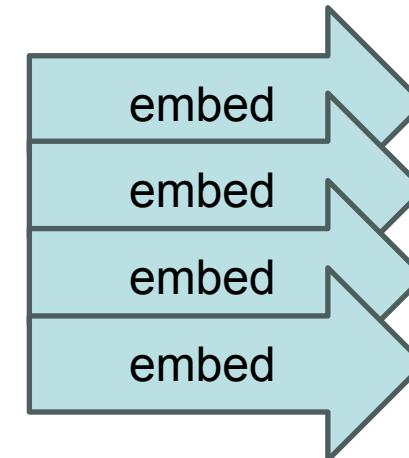
- Input: some text

- “The”
- “ dog”
- “ chased”
- “ the”



one-hot

[791]  
[5679]  
[62920]  
[279]

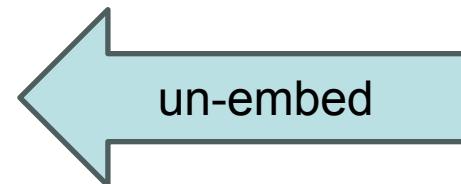


- Output: more text

- “ ball”

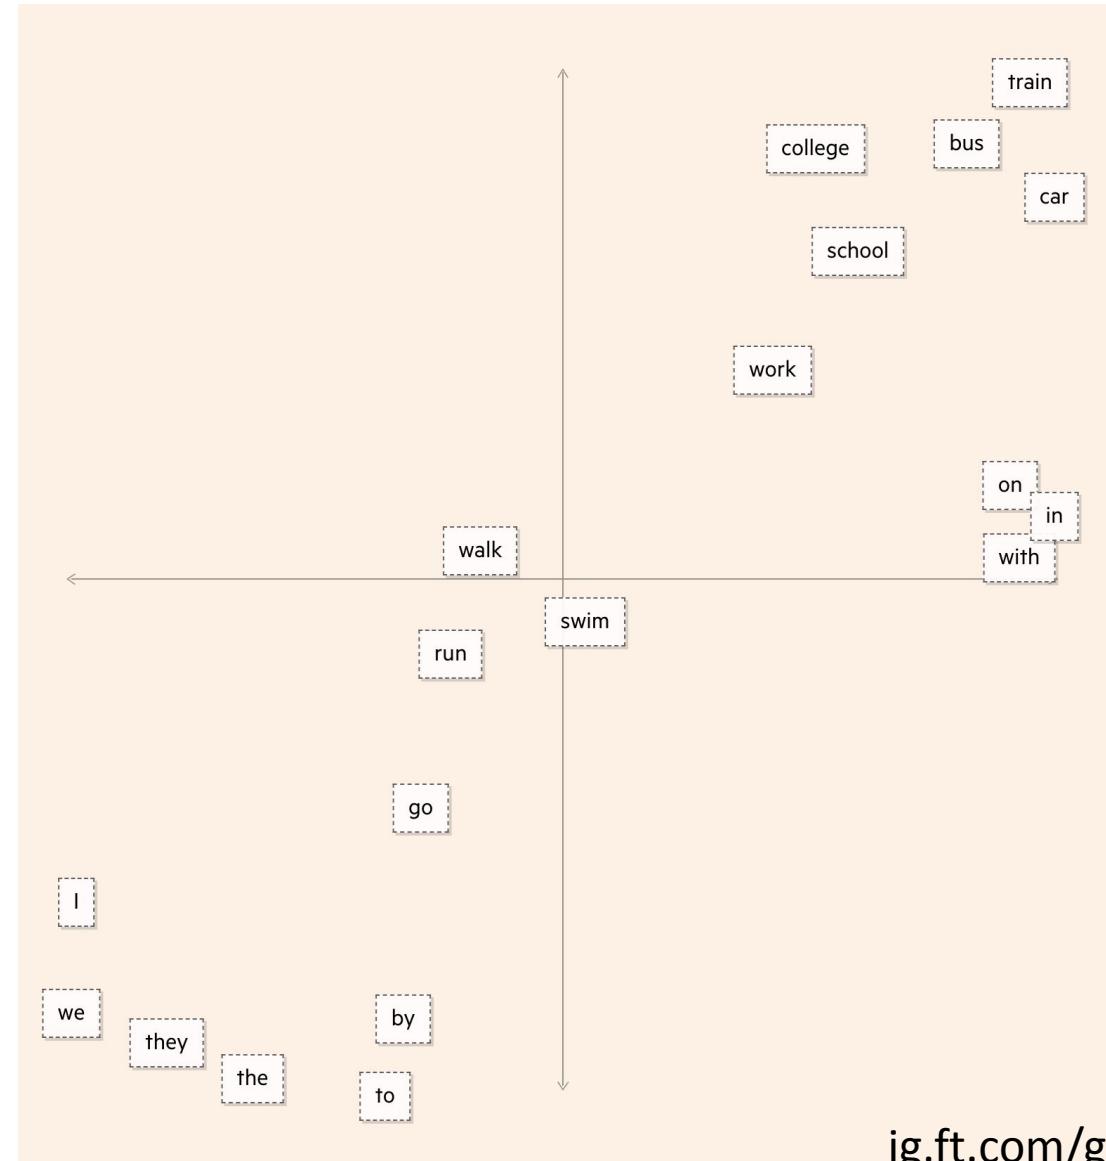


[5041]



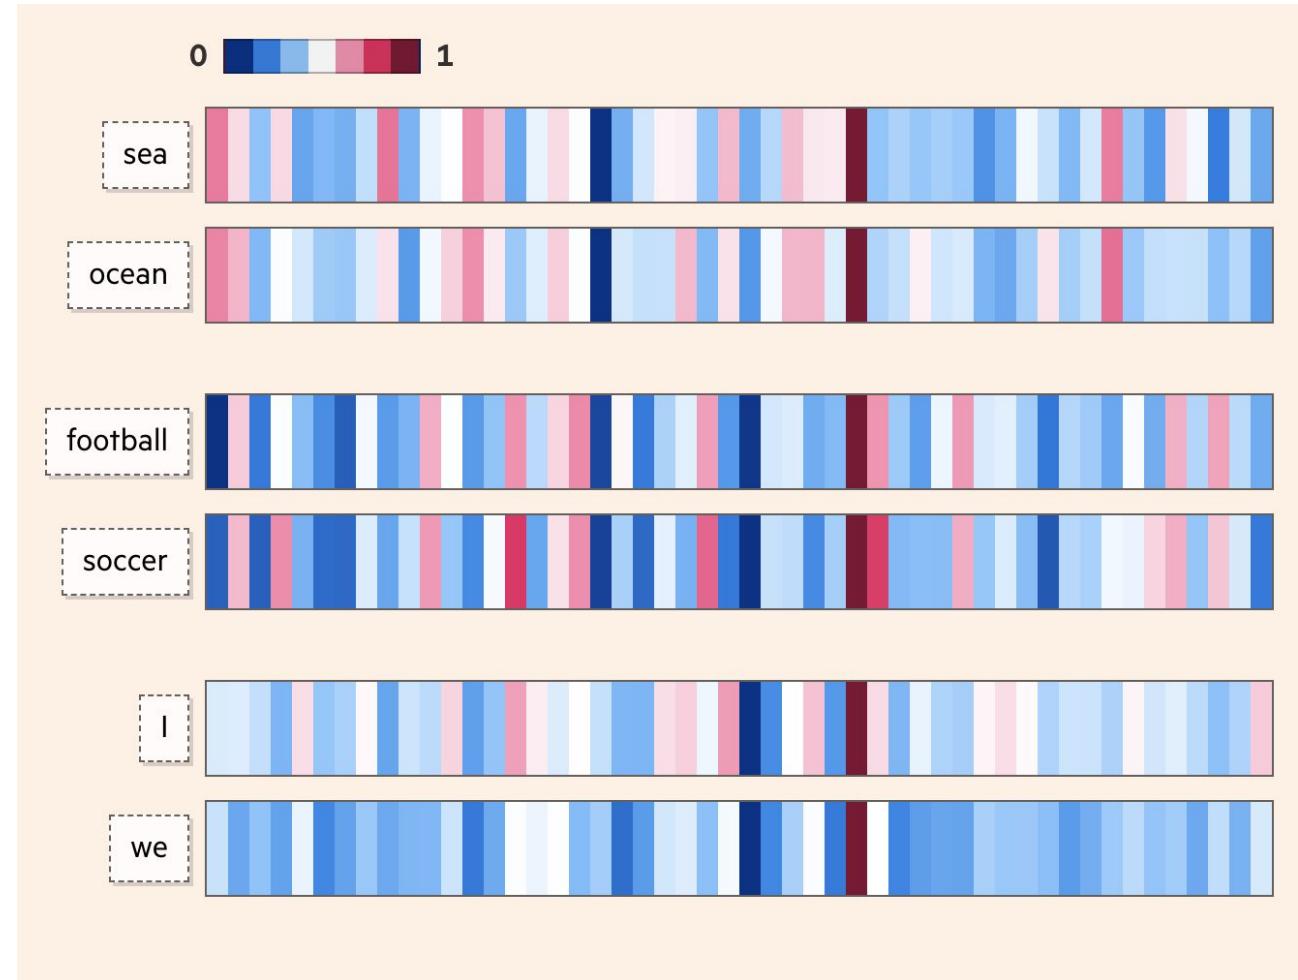
# What do word embeddings look like?

- Words cluster by similarity:



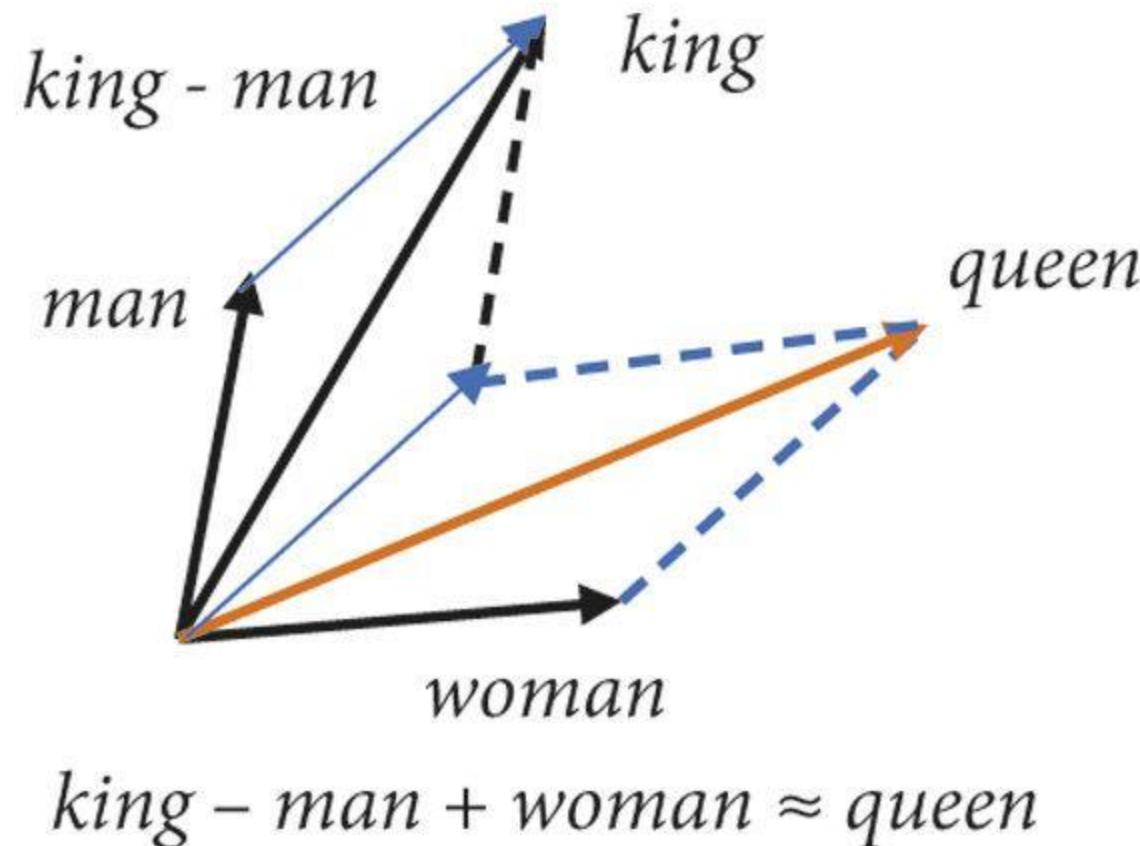
# What do word embeddings look like?

- Features learned in language models:



# What do word embeddings look like?

- Signs of sensible algebra in embedding space:



[Efficient estimation of word representations in vector space, Mikolov et al, 2013]

# What do word embeddings look like?

---

**Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings**

Tolga Bolukbasi, Kai-Wei Chang, James Zou, Venkatesh Saligrama, Adam Kalai

# Aside: interactive explainer of modern language models

---

[ig.ft.com/generative-ai](https://ig.ft.com/generative-ai)

Artificial Intelligence

## Generative AI exists because of the transformer

This is how it works

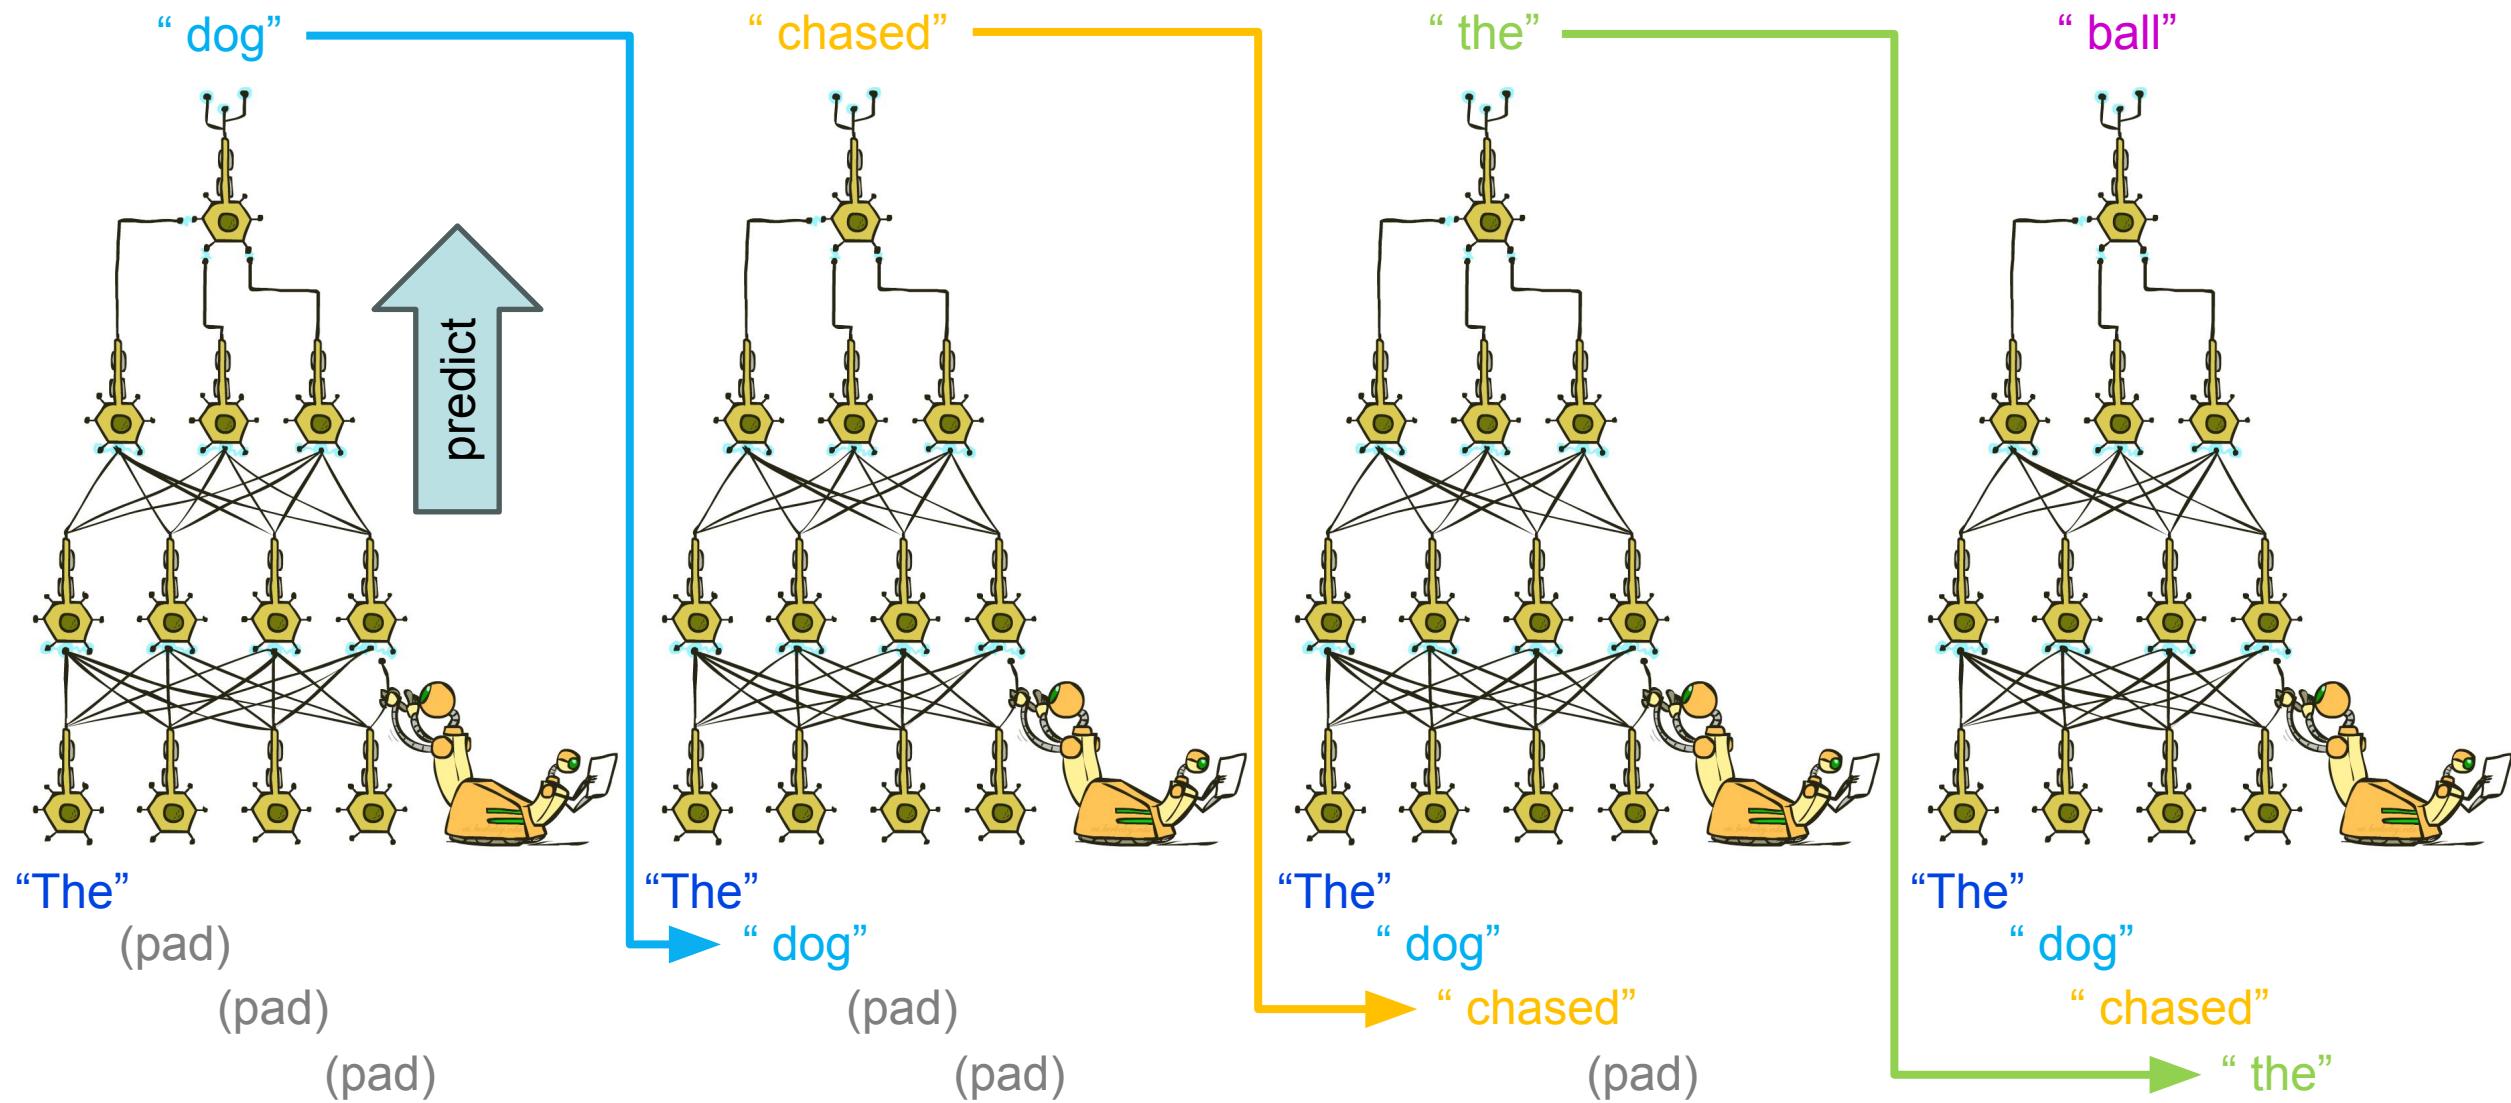
By Visual Storytelling Team and Madhumita Murgia in London SEPTEMBER 11 2023

# Large Language Models

---

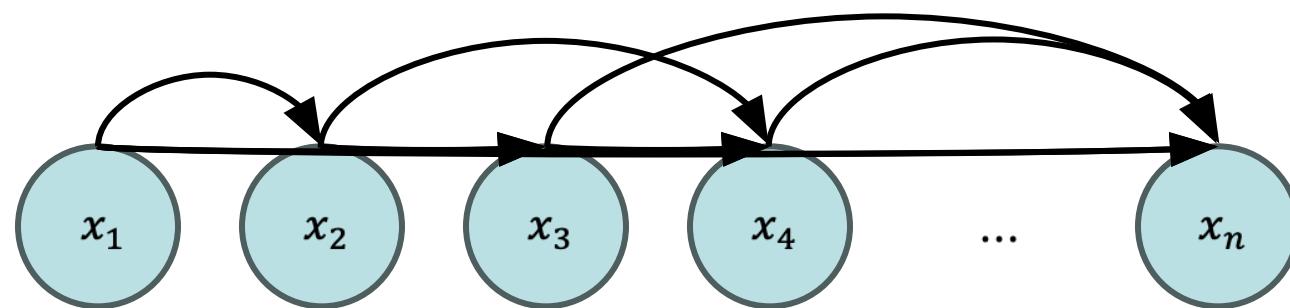
- ~~Feature engineering~~
  - Text tokenization
  - Word embeddings
- Deep neural networks
  - Autoregressive models
  - Self-attention mechanisms
  - Transformer architectures
- Multi-class classification
- Supervised learning
  - Self-supervised learning
  - Instruction tuning
- Reinforcement learning
  - ... from human feedback (RLHF)

# Autoregressive Models



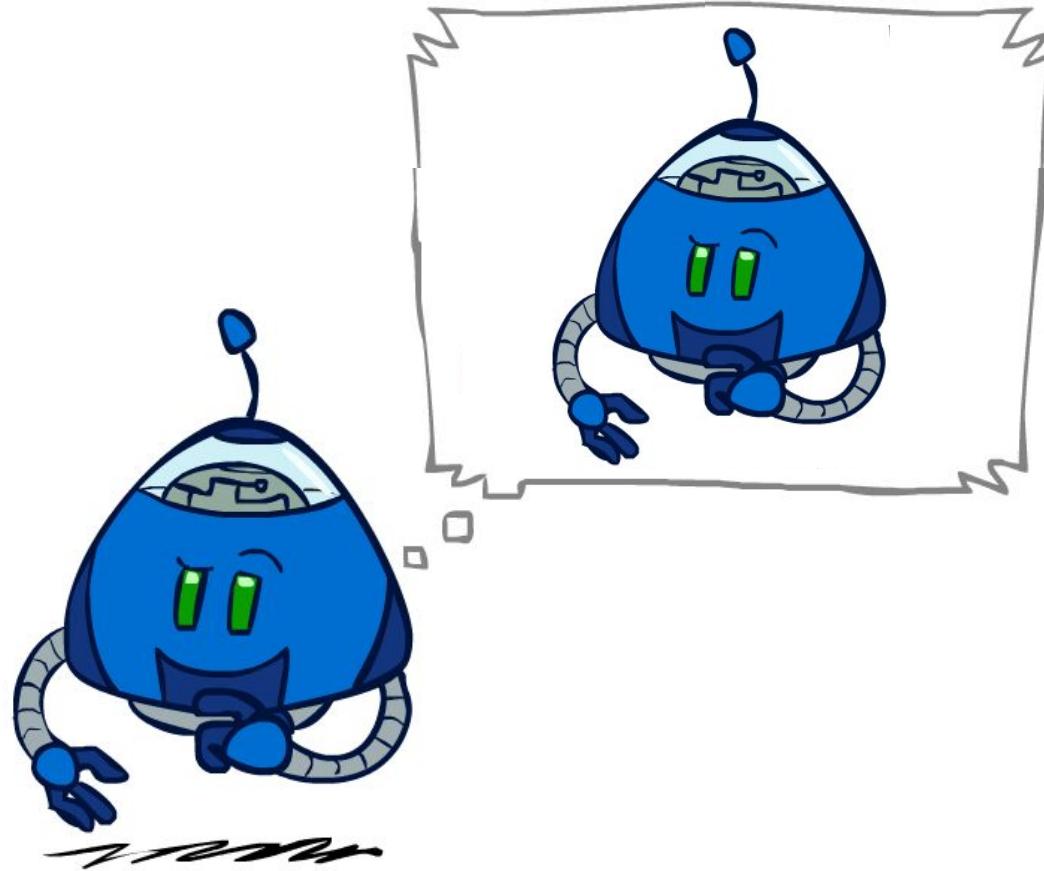
# Autoregressive Models

- Predict output one piece at a time (e.g. word, token, pixel, etc.)
- Concatenate: input + output
- Feed result back in as new input
- Repeat

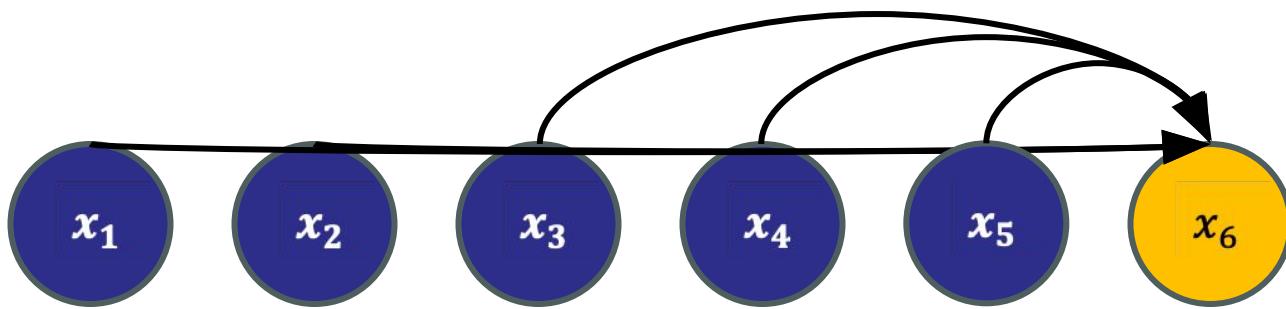


# Self-Attention Mechanisms

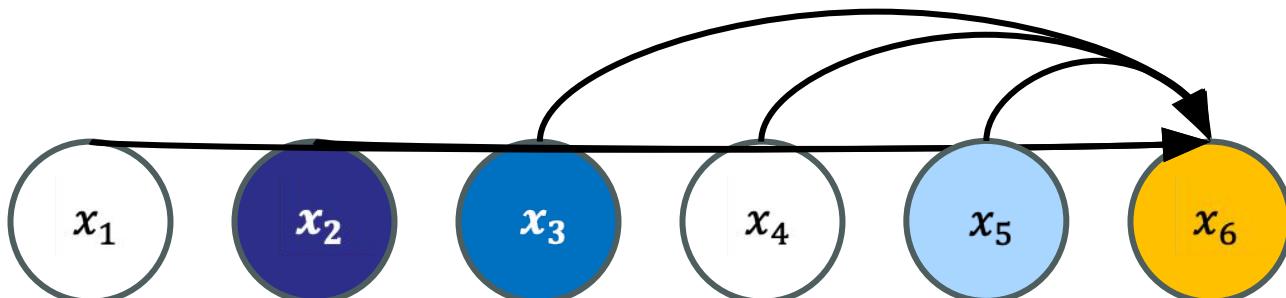
---



# Self-Attention Mechanisms



- Instead of conditioning on *all* input tokens equally...



- Pay more attention to relevant tokens!

# Self-Attention Mechanisms



*output*

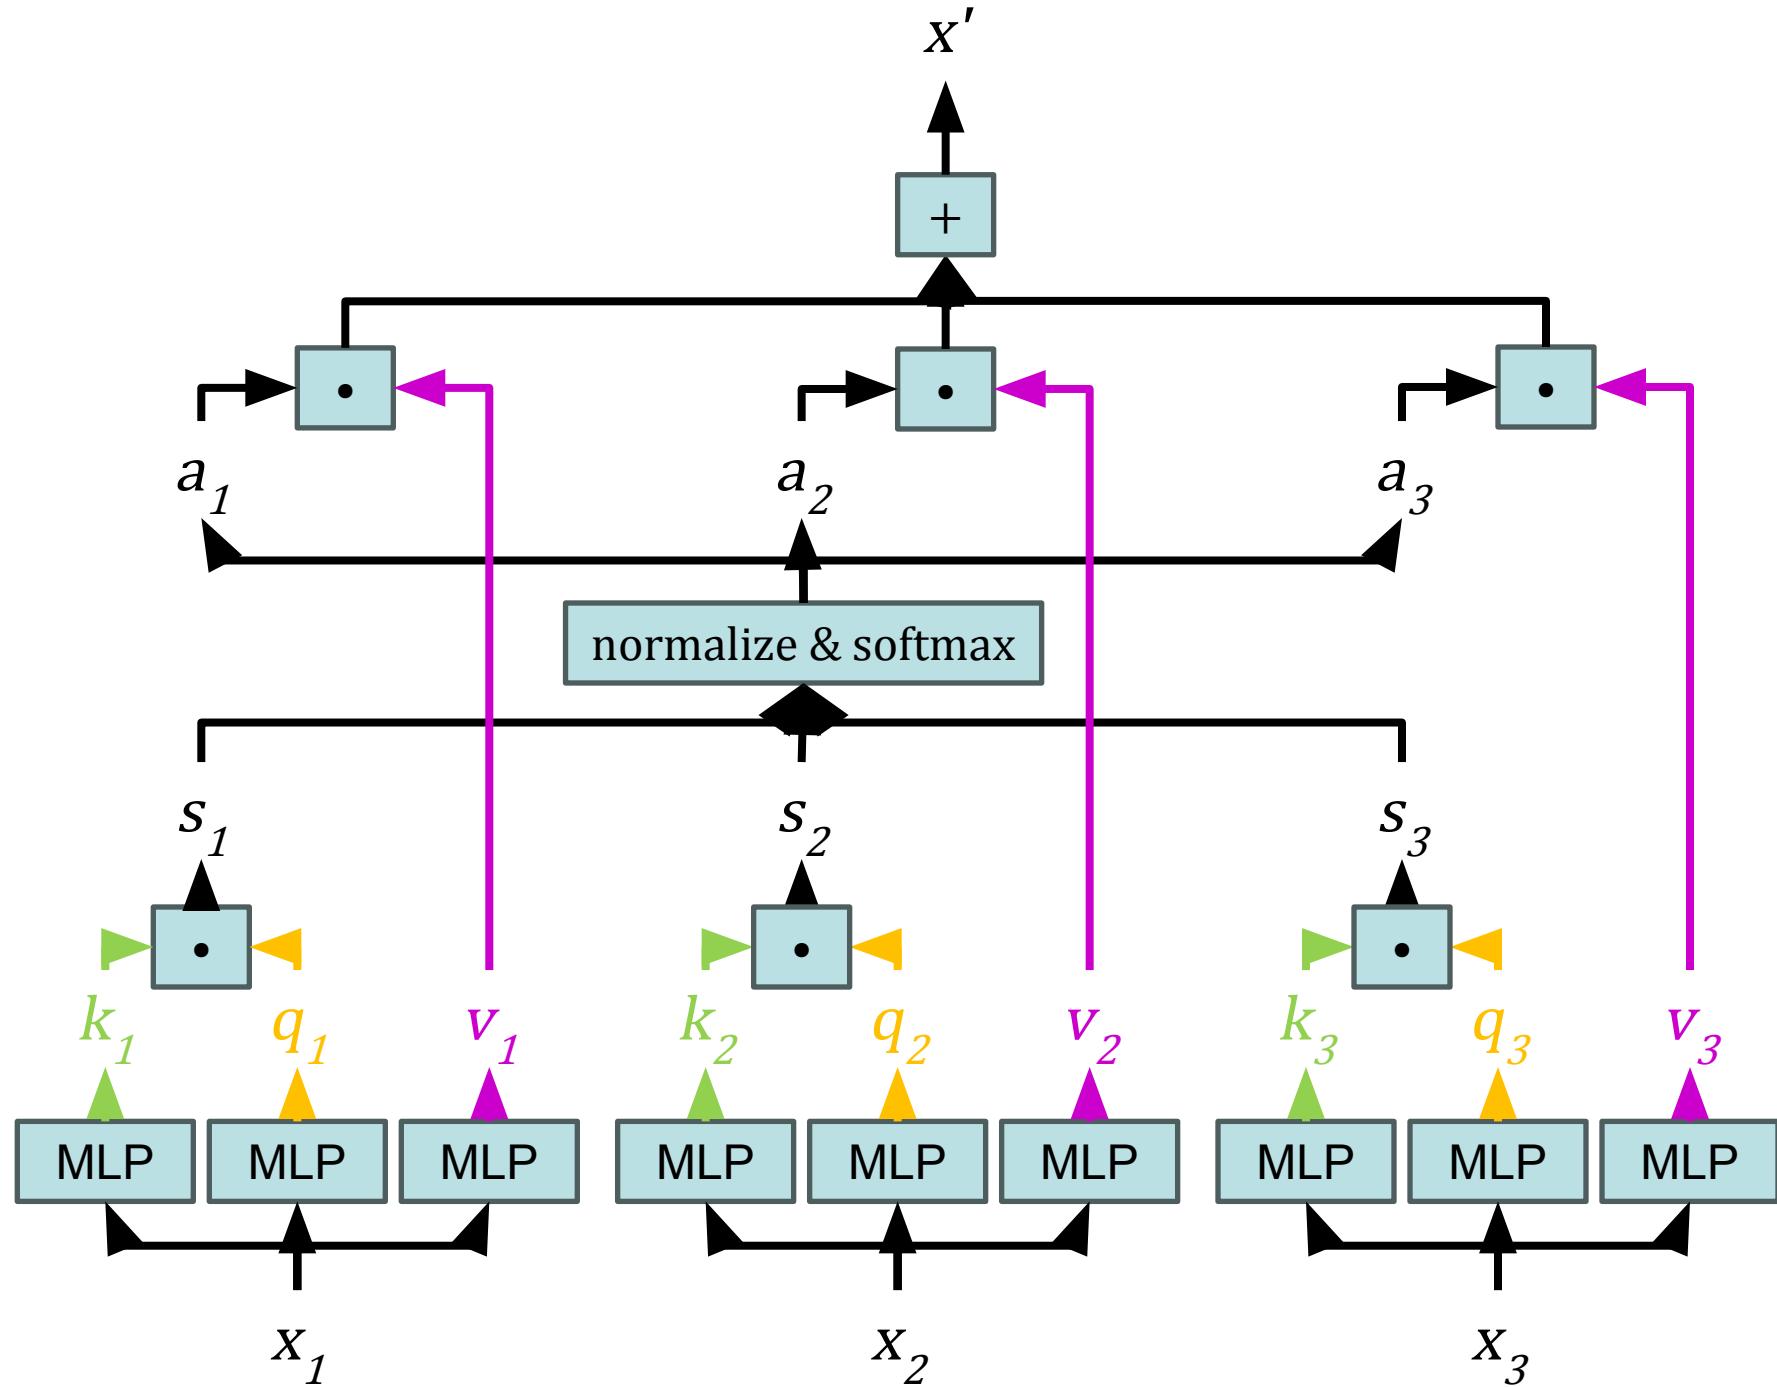
*attention weight*

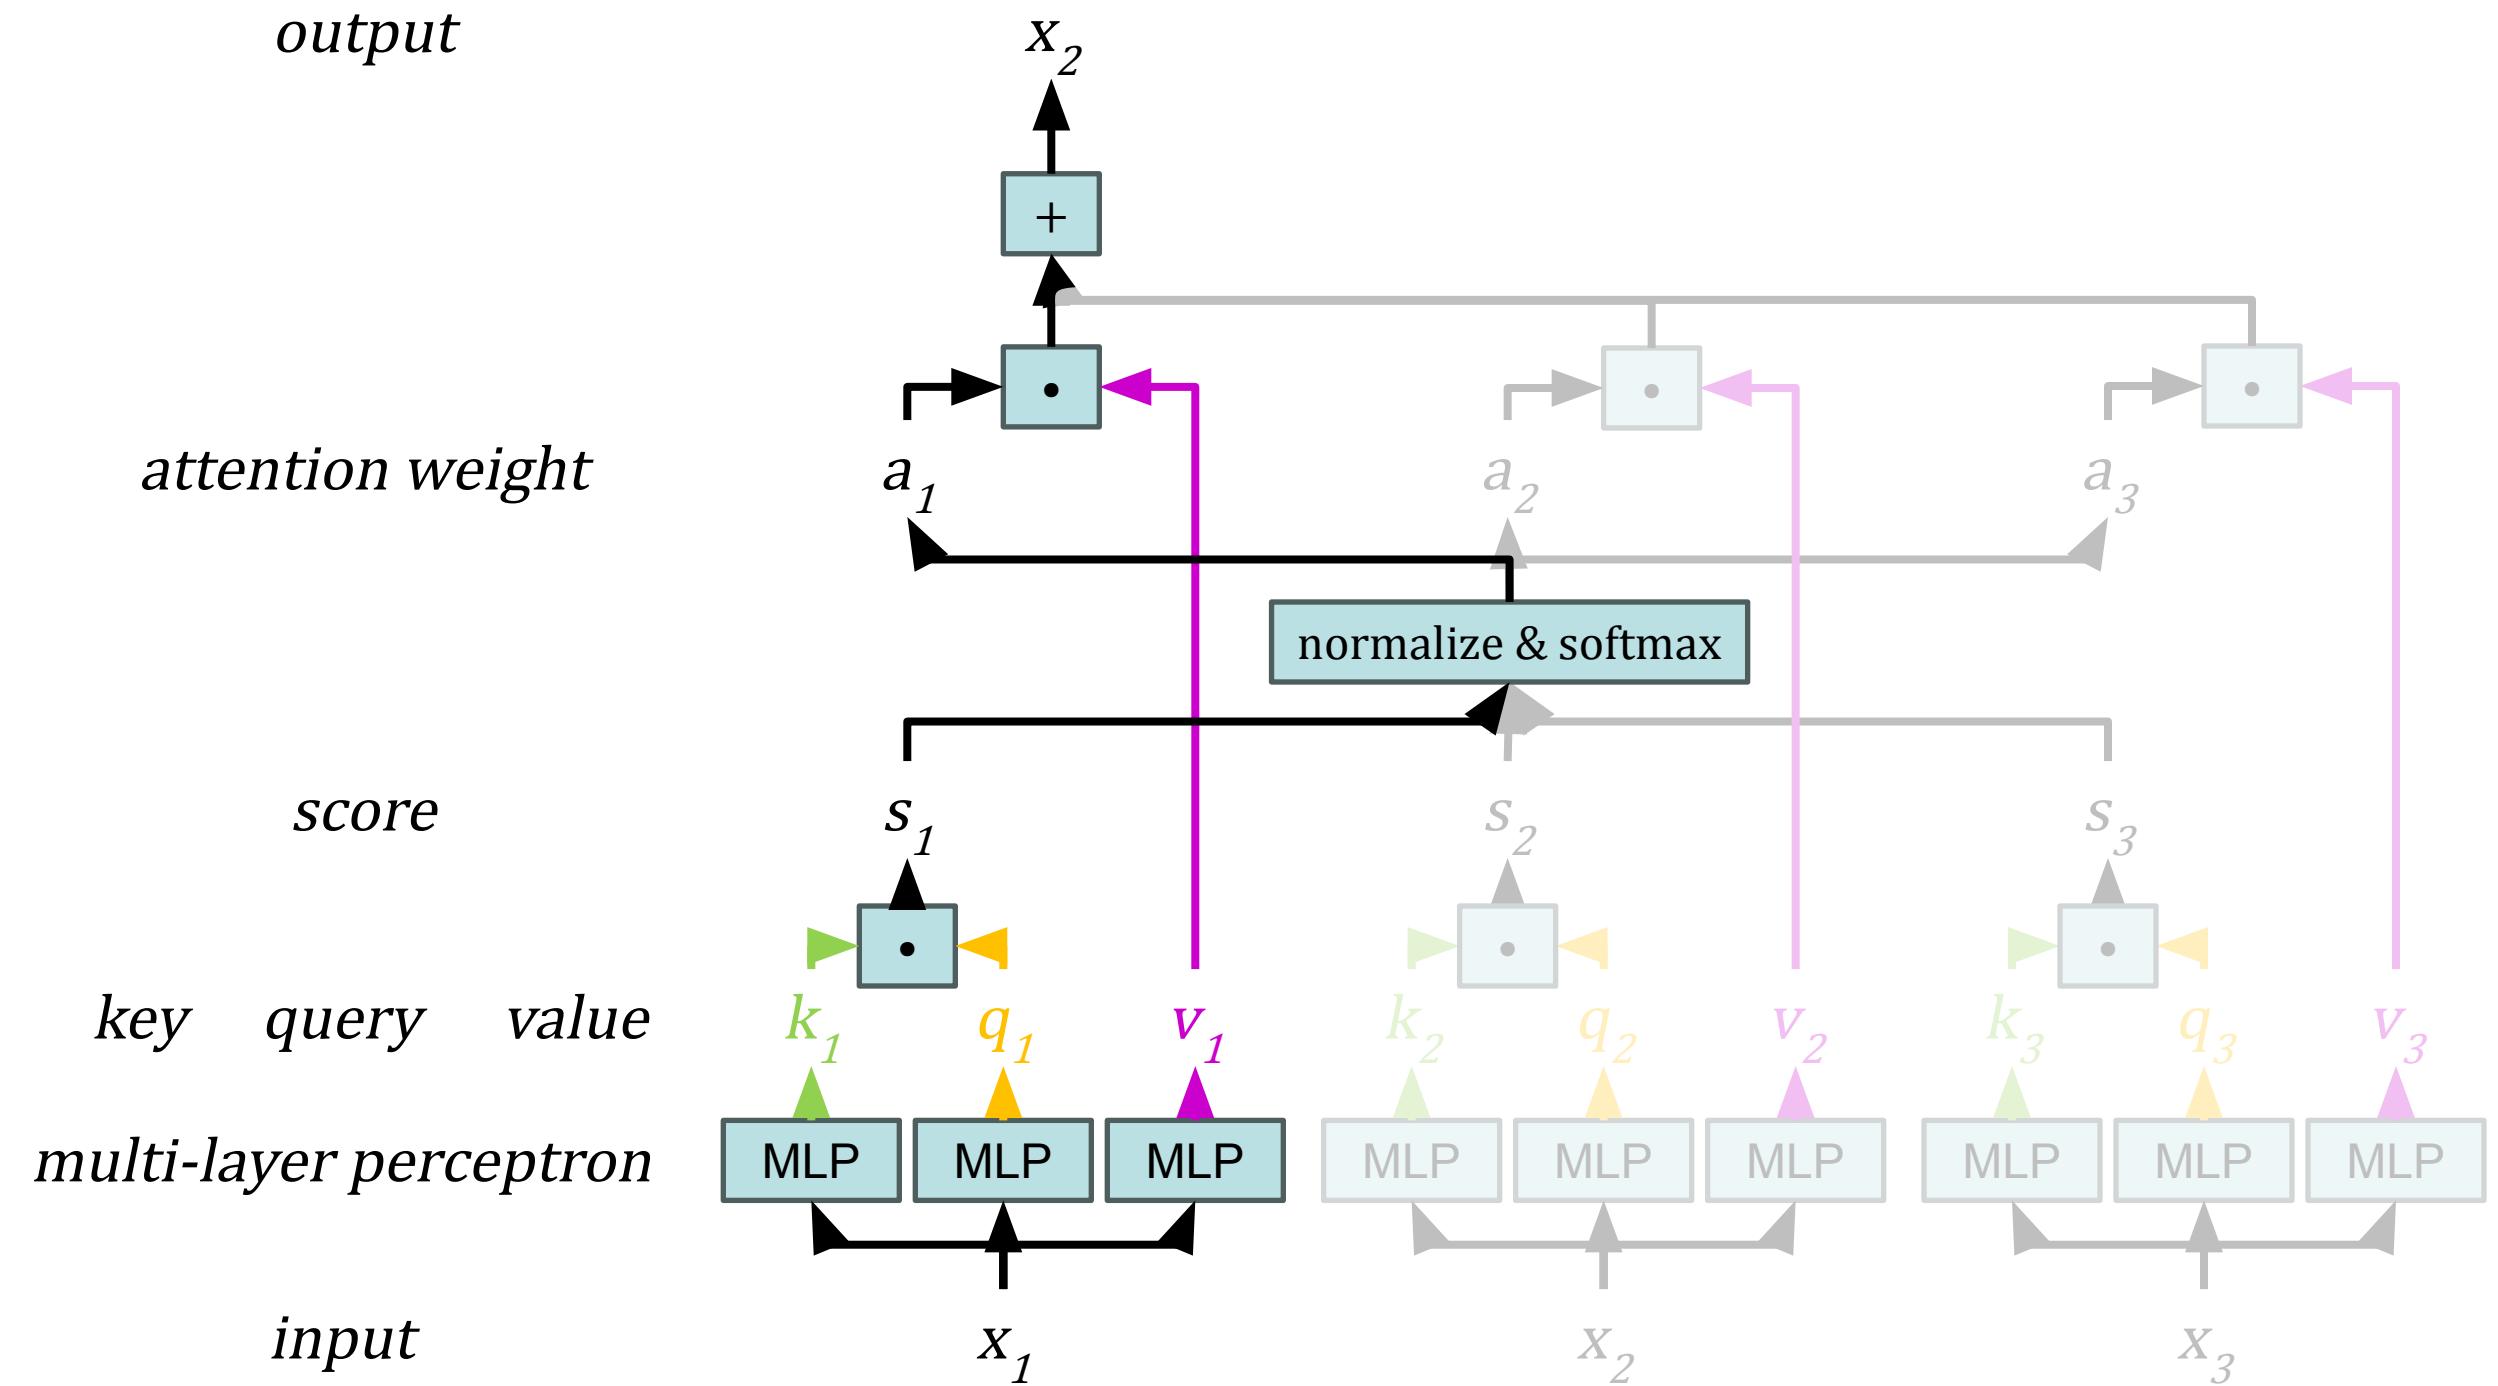
*score*

*key    query    value*

*multi-layer perceptron*

*input*





*output*

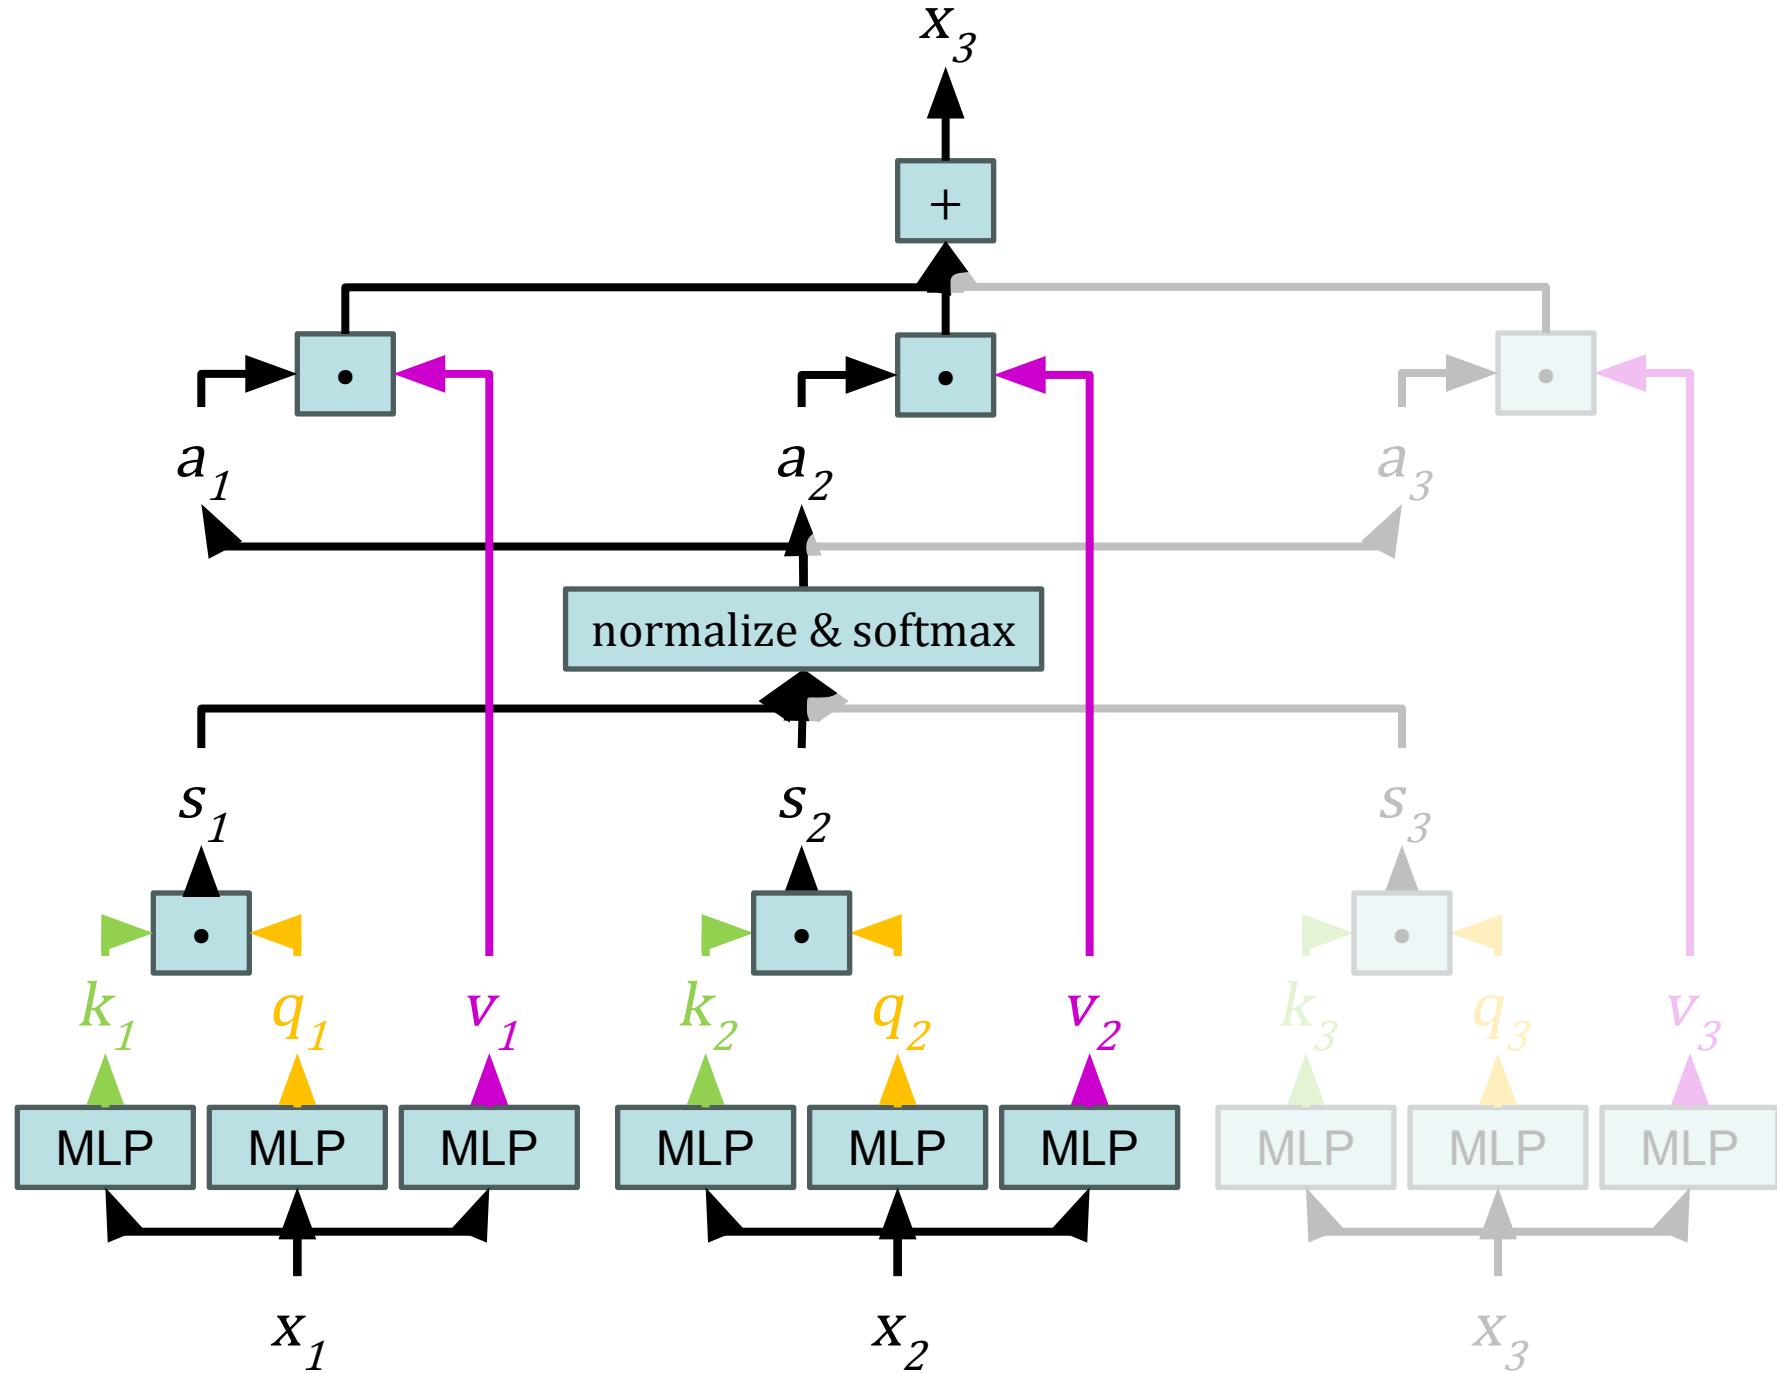
*attention weight*

*score*

*key    query    value*

*multi-layer perceptron*

*input*



*output*

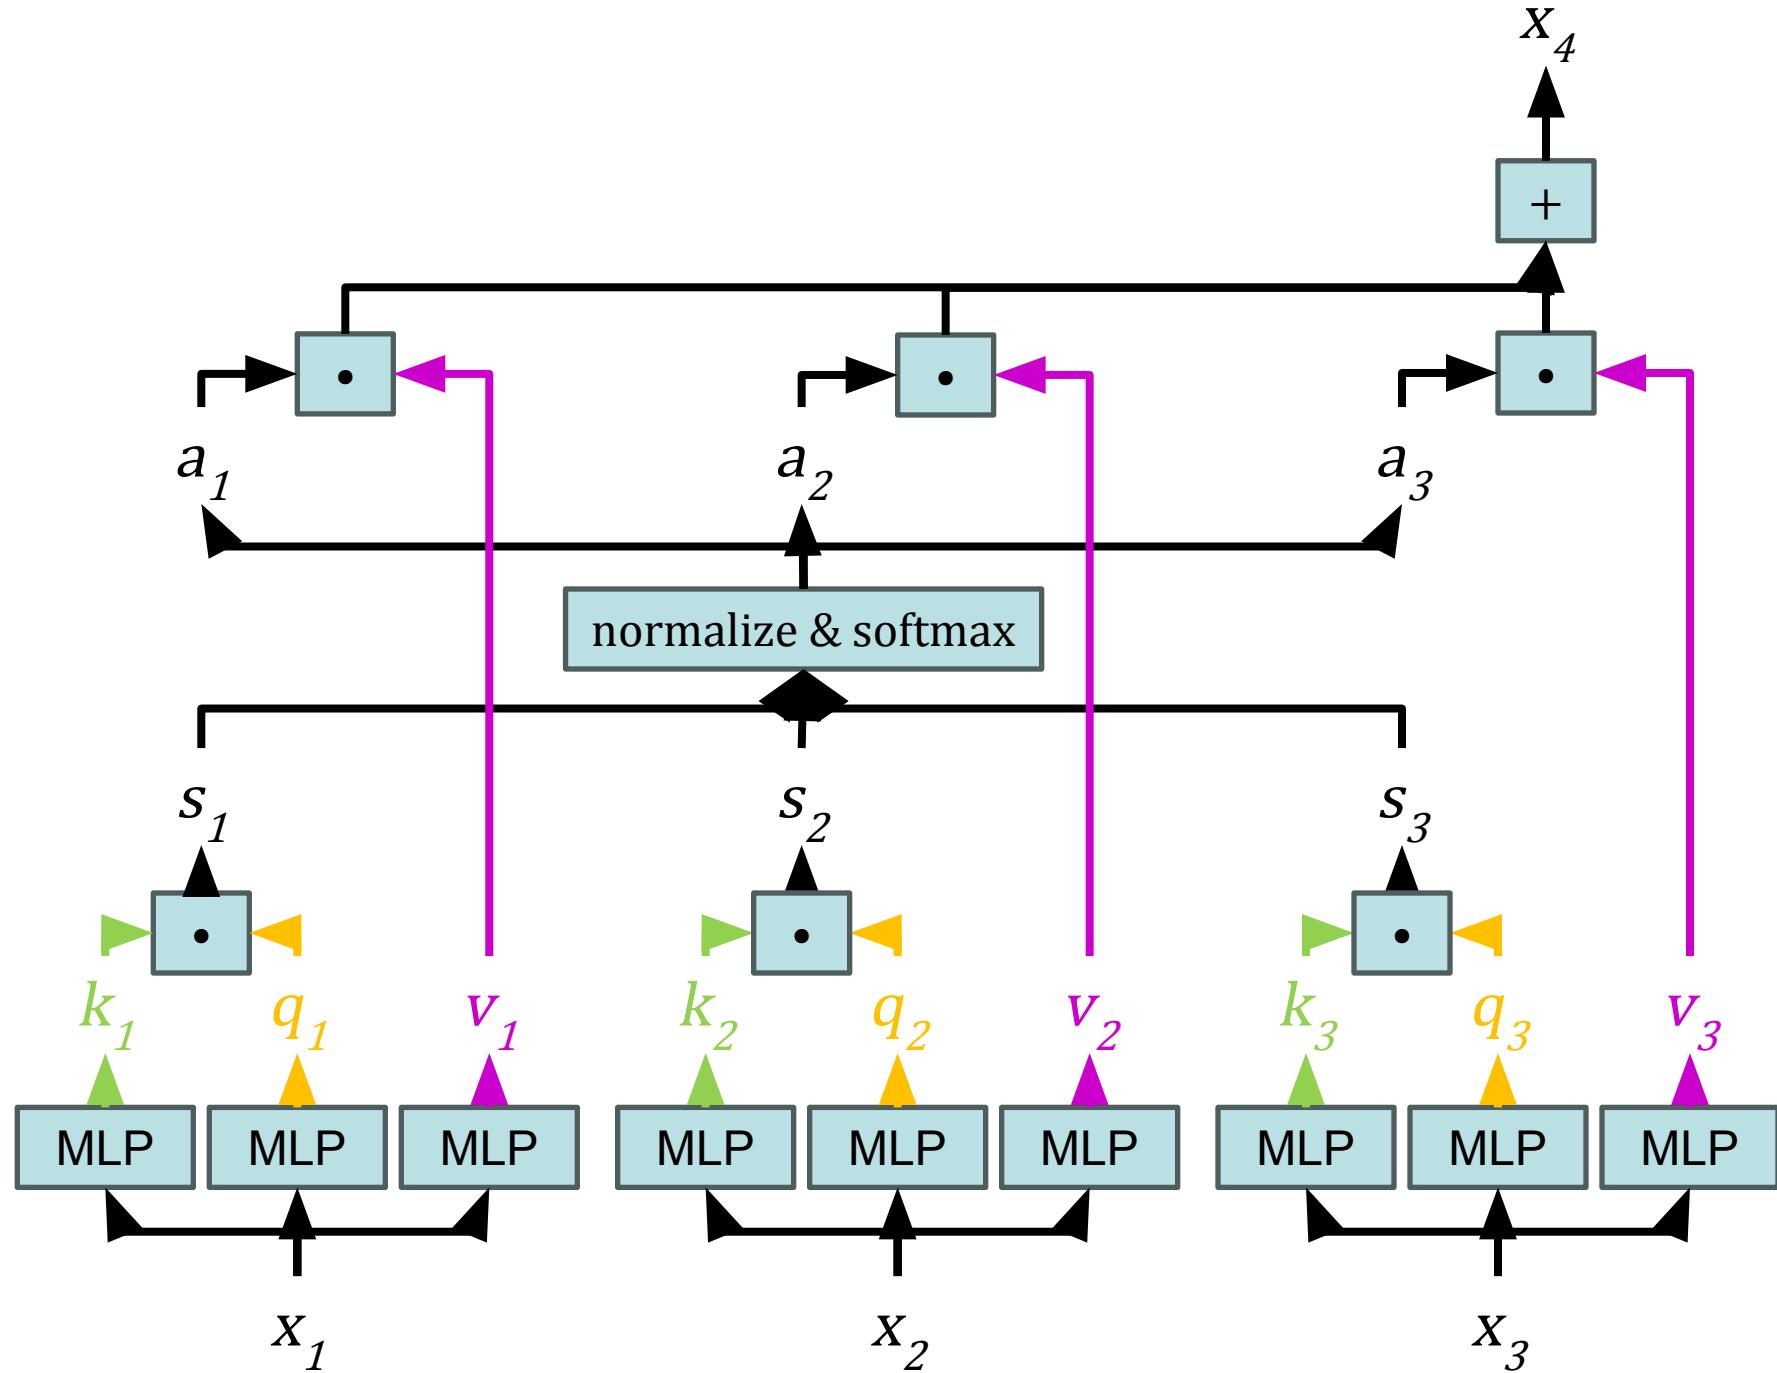
*attention weight*

*score*

*key    query    value*

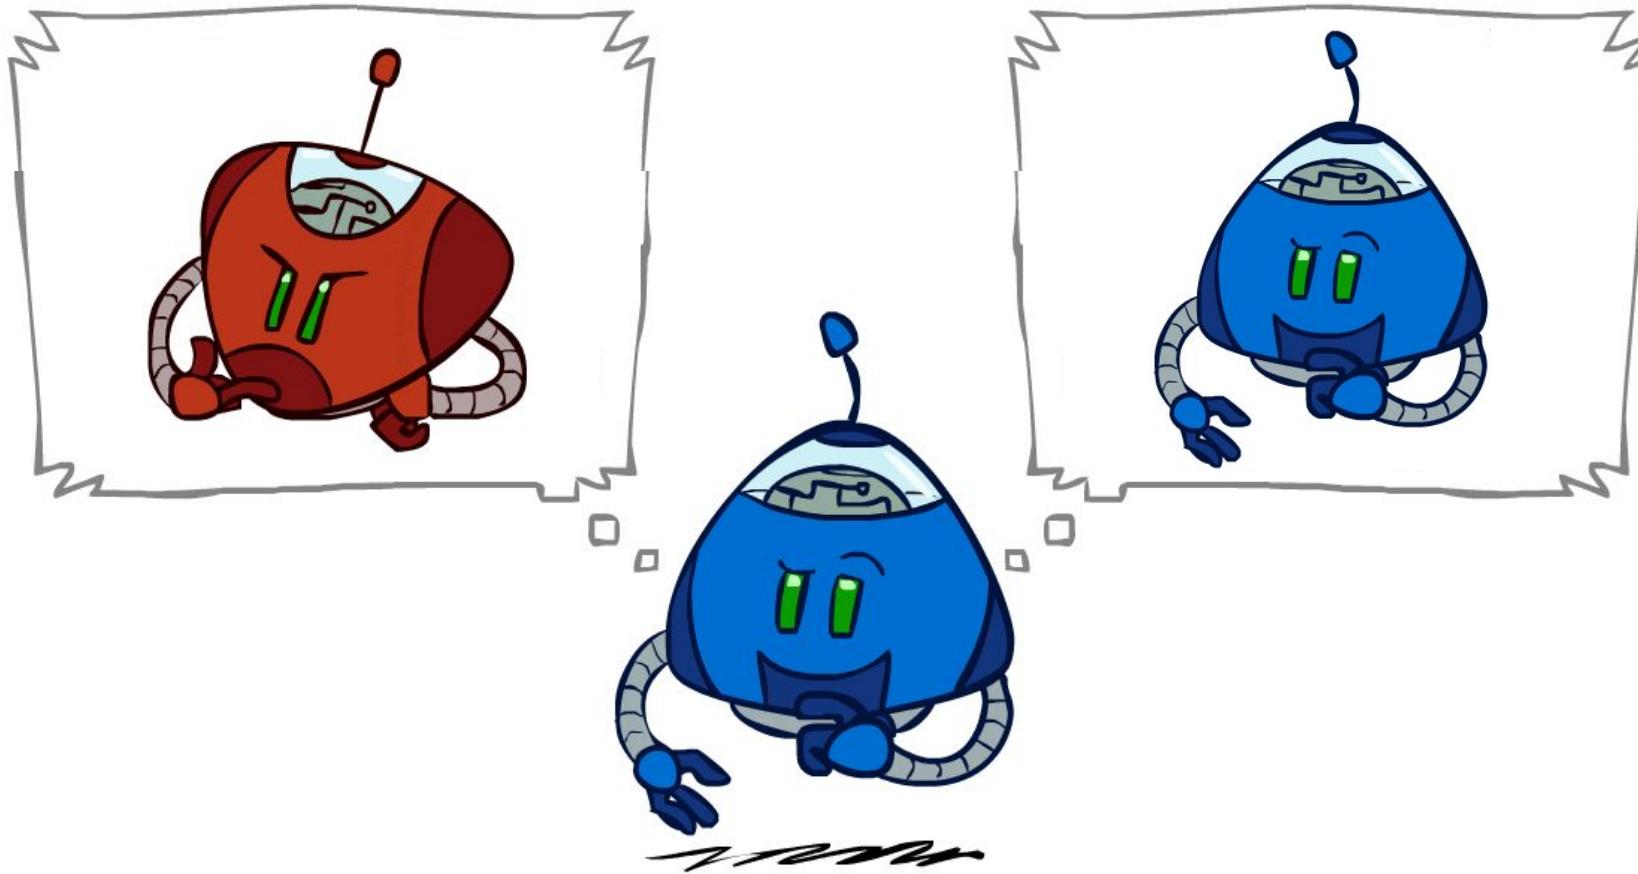
*multi-layer perceptron*

*input*



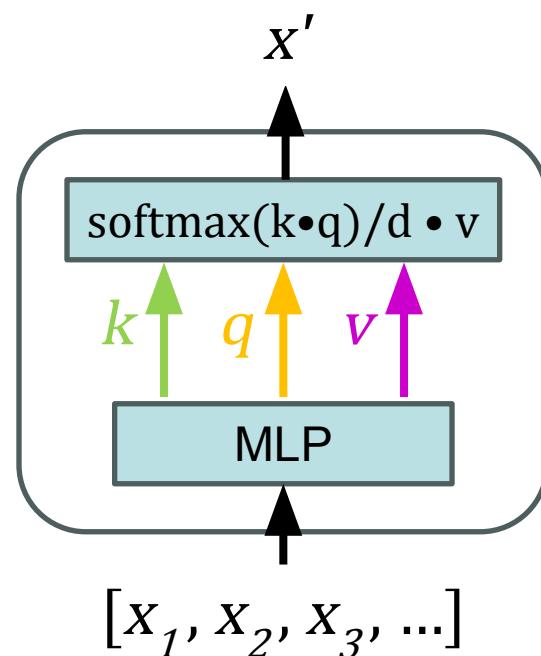
# Multi-Headed Attention

---

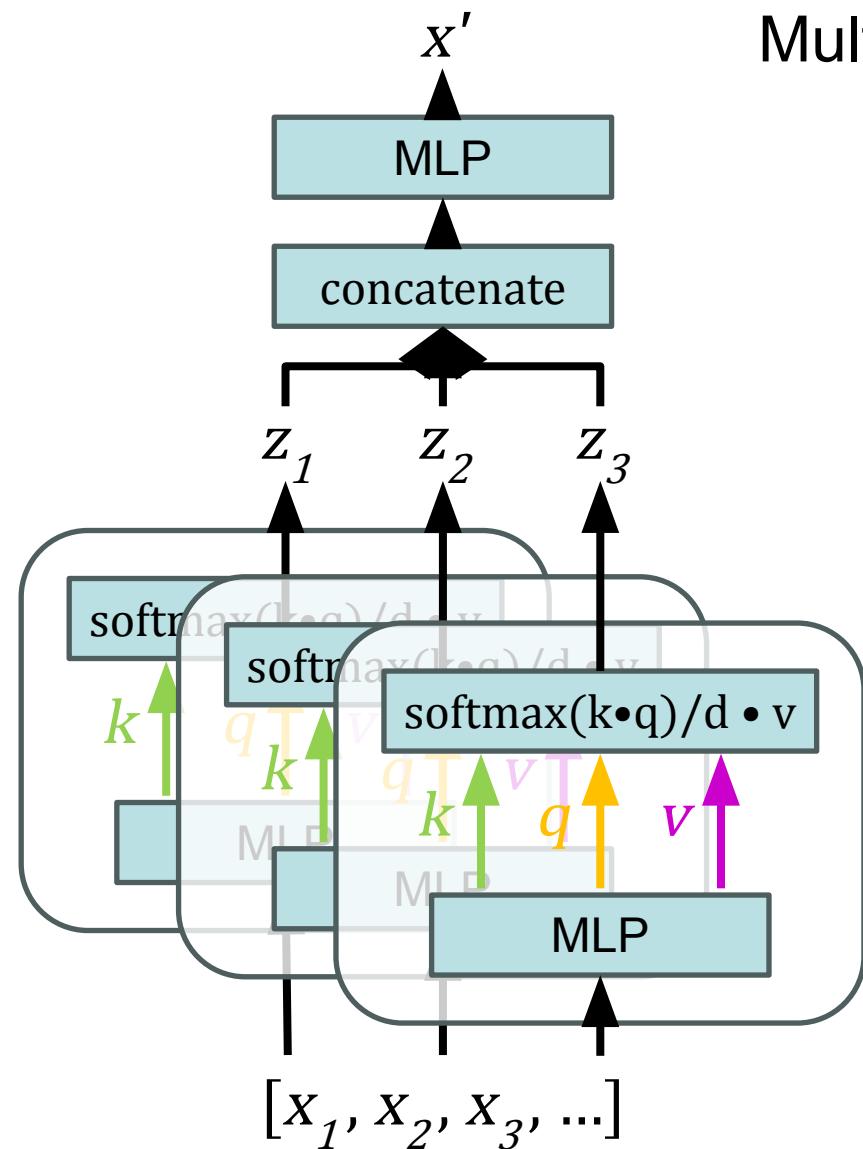


# Multi-Headed Attention

Single-headed

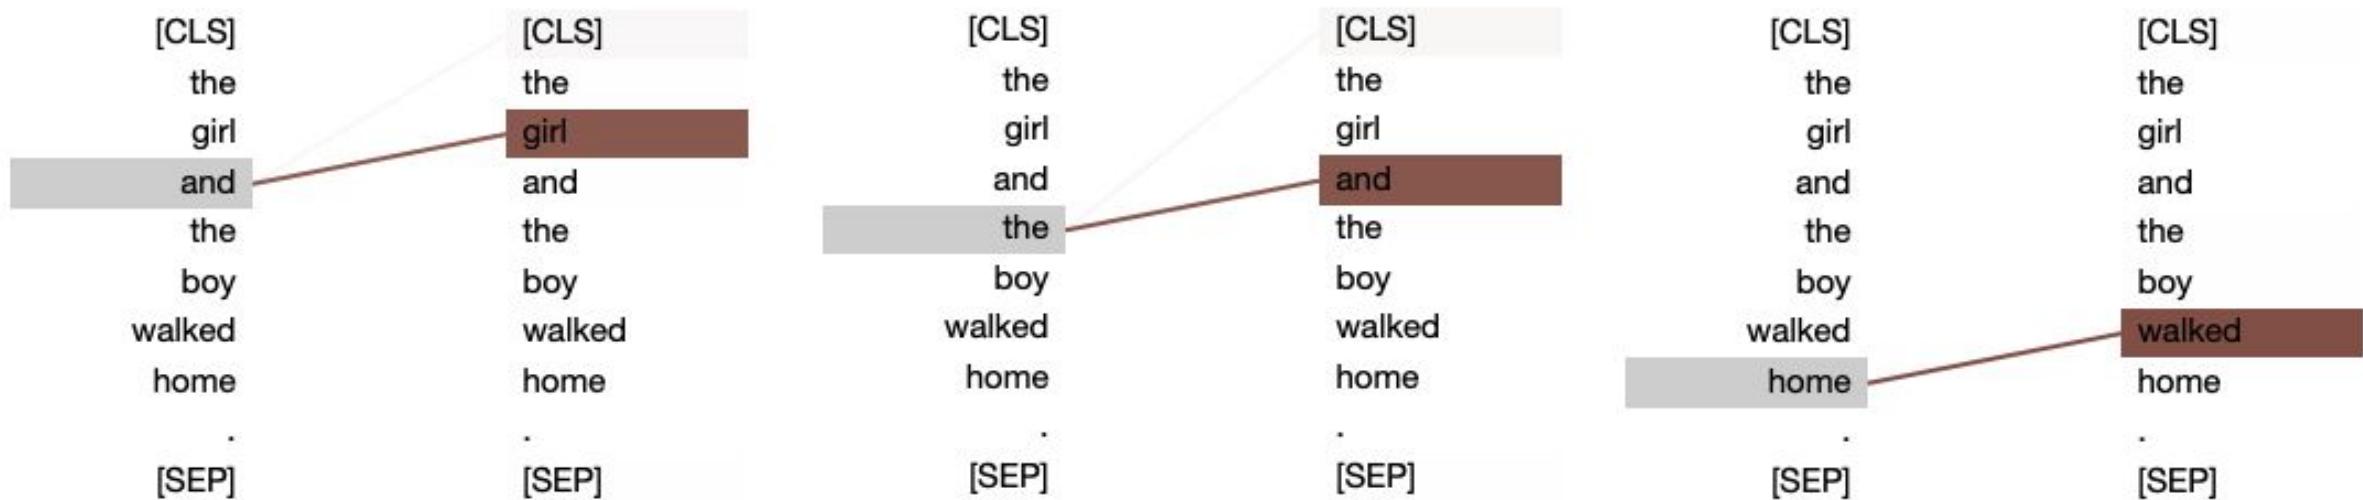


Multi-headed



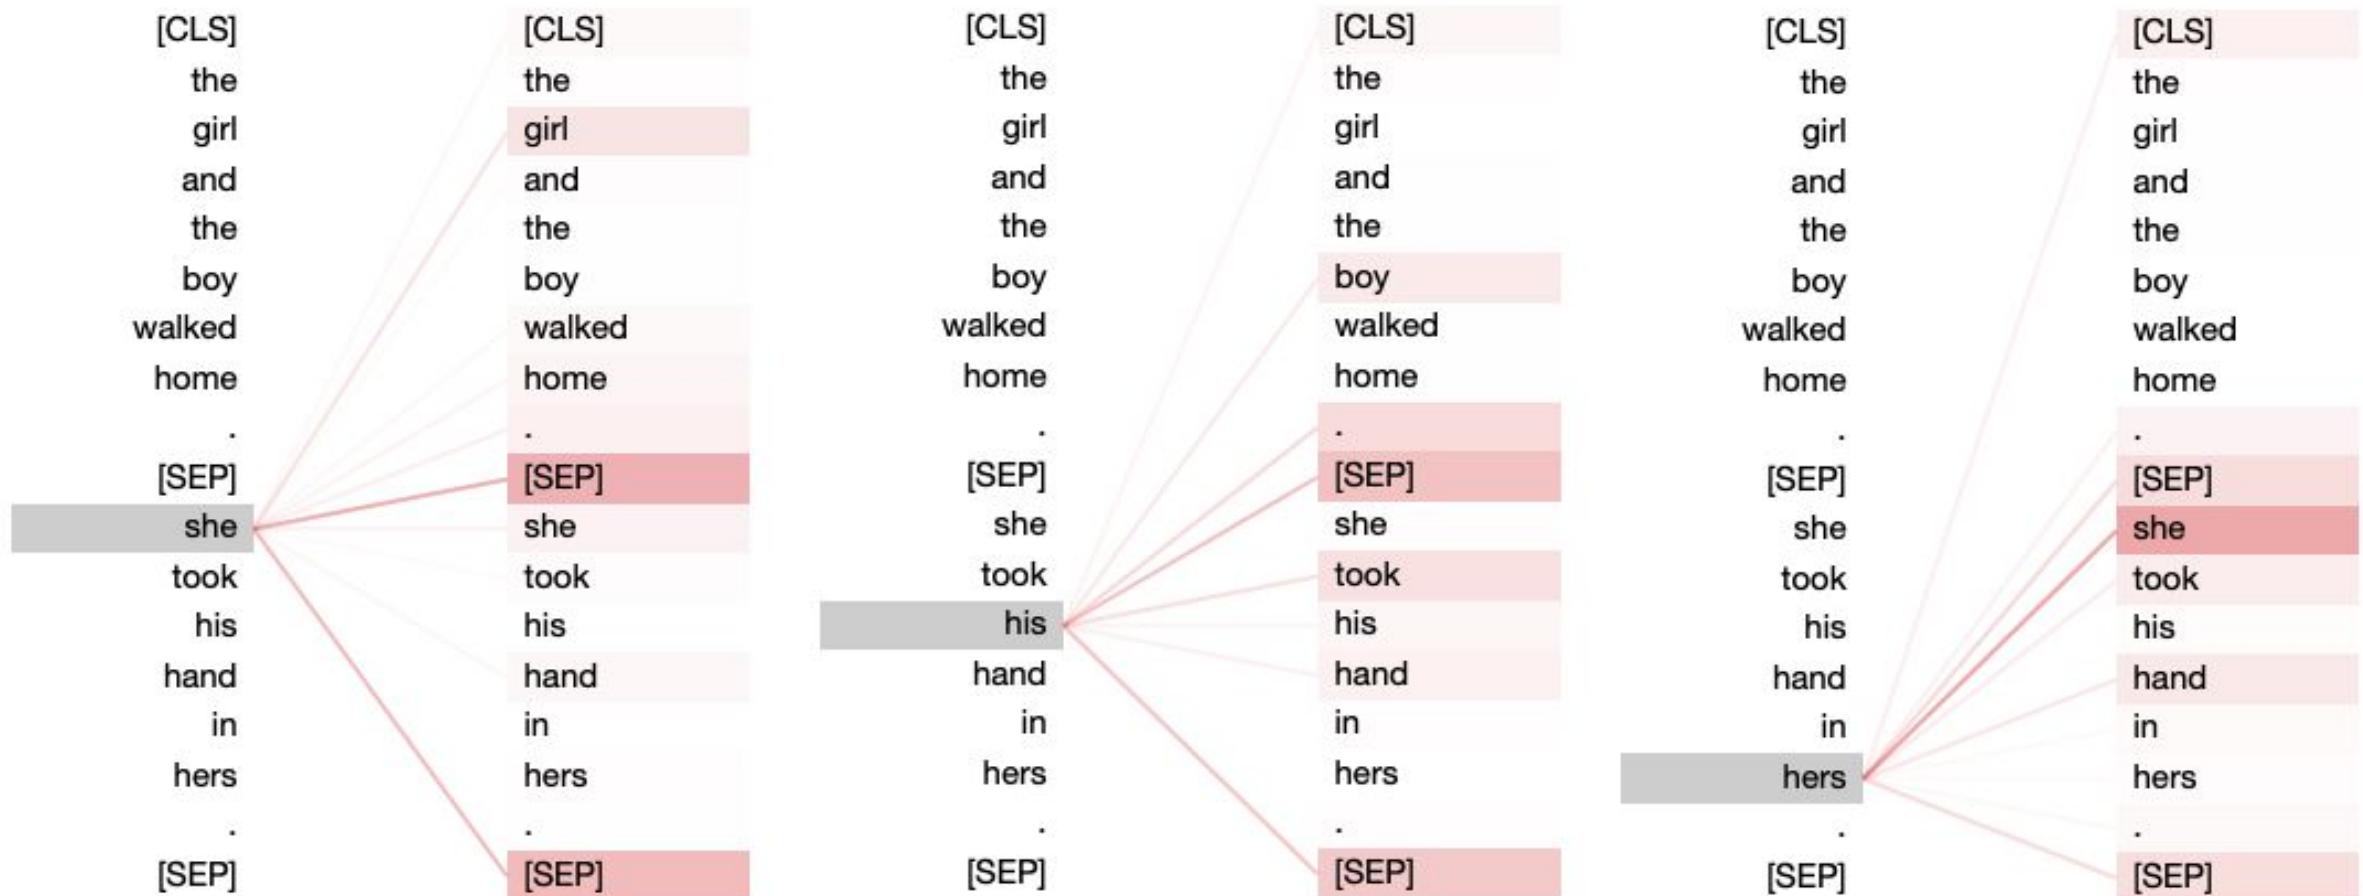
# Multi-Headed Attention

Head 6: previous word



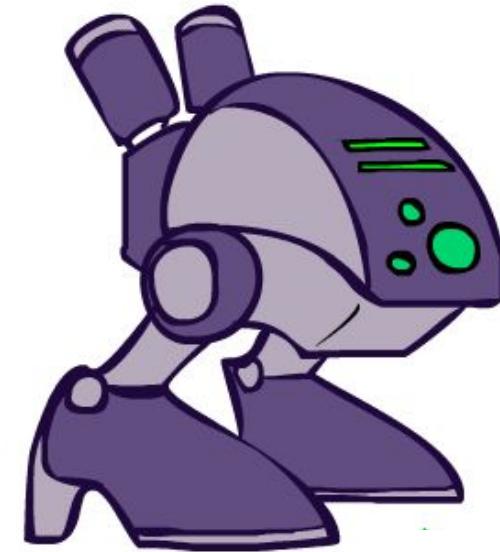
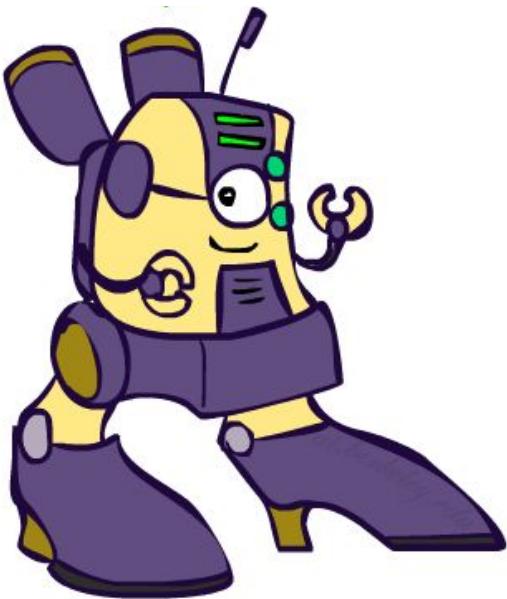
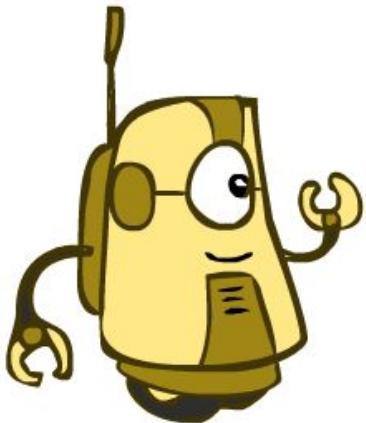
# Multi-Headed Attention

## Head 4: pronoun references

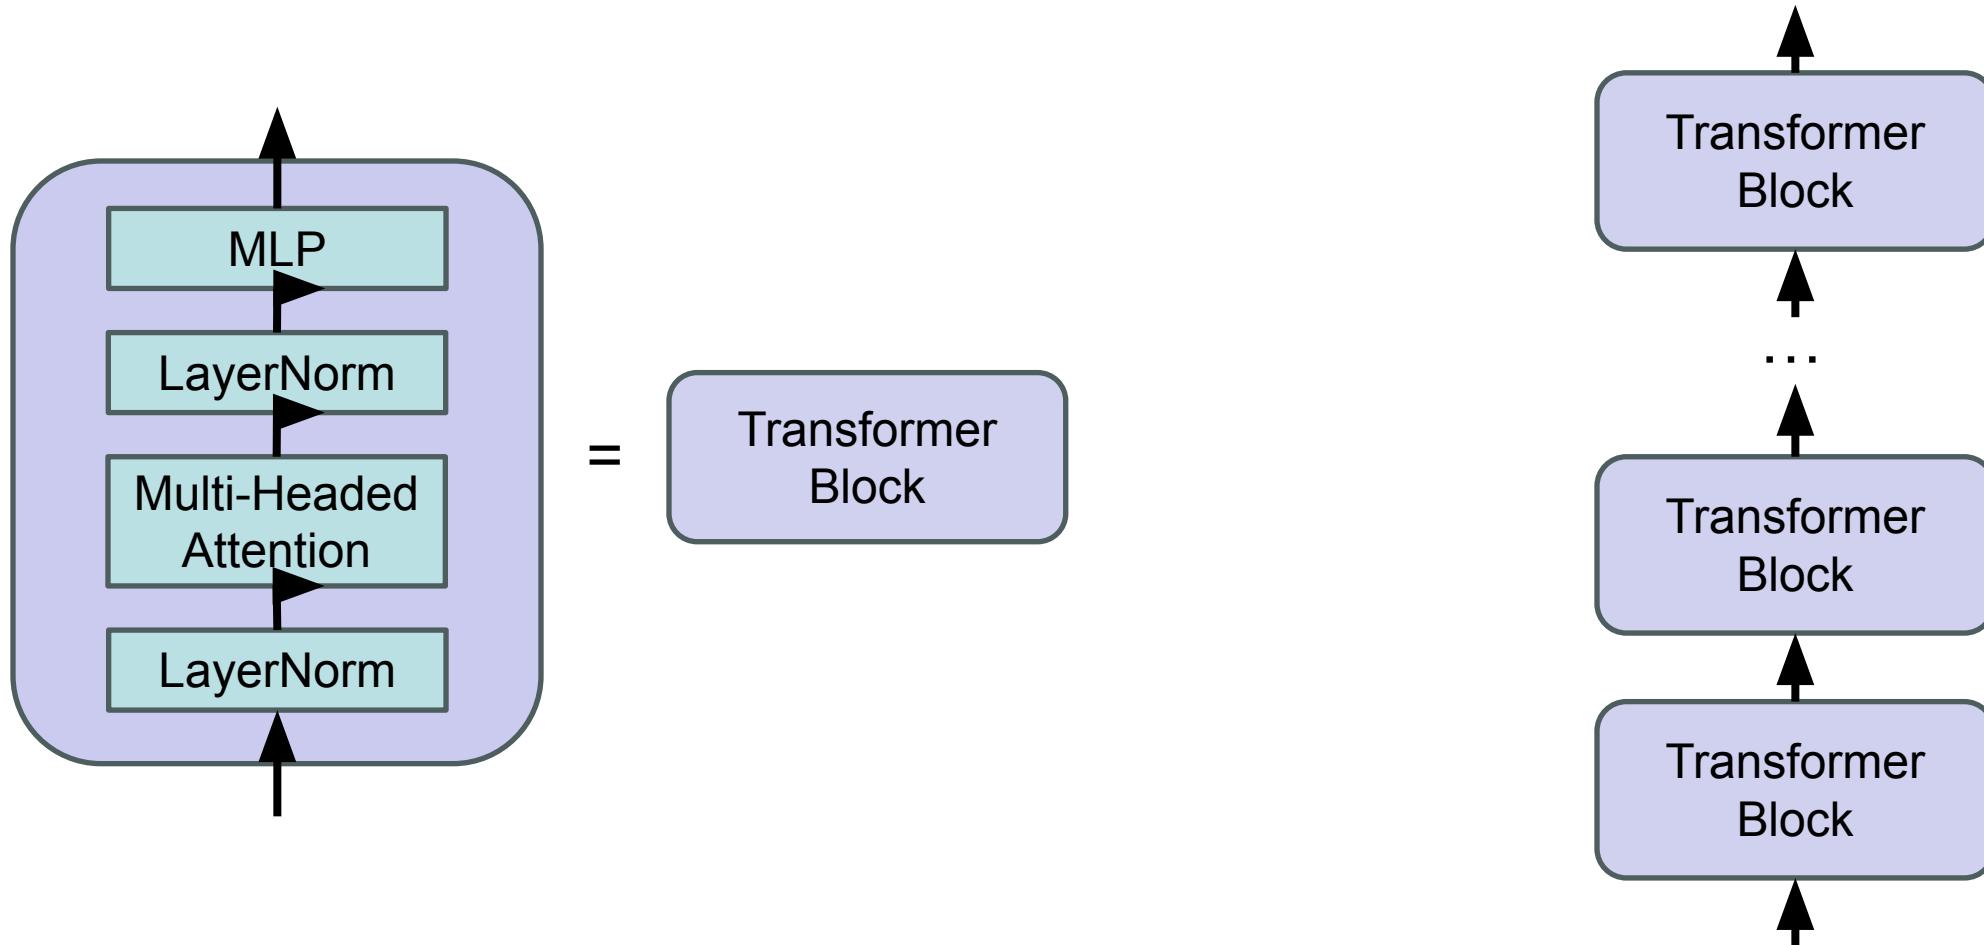


# Transformer Architecture

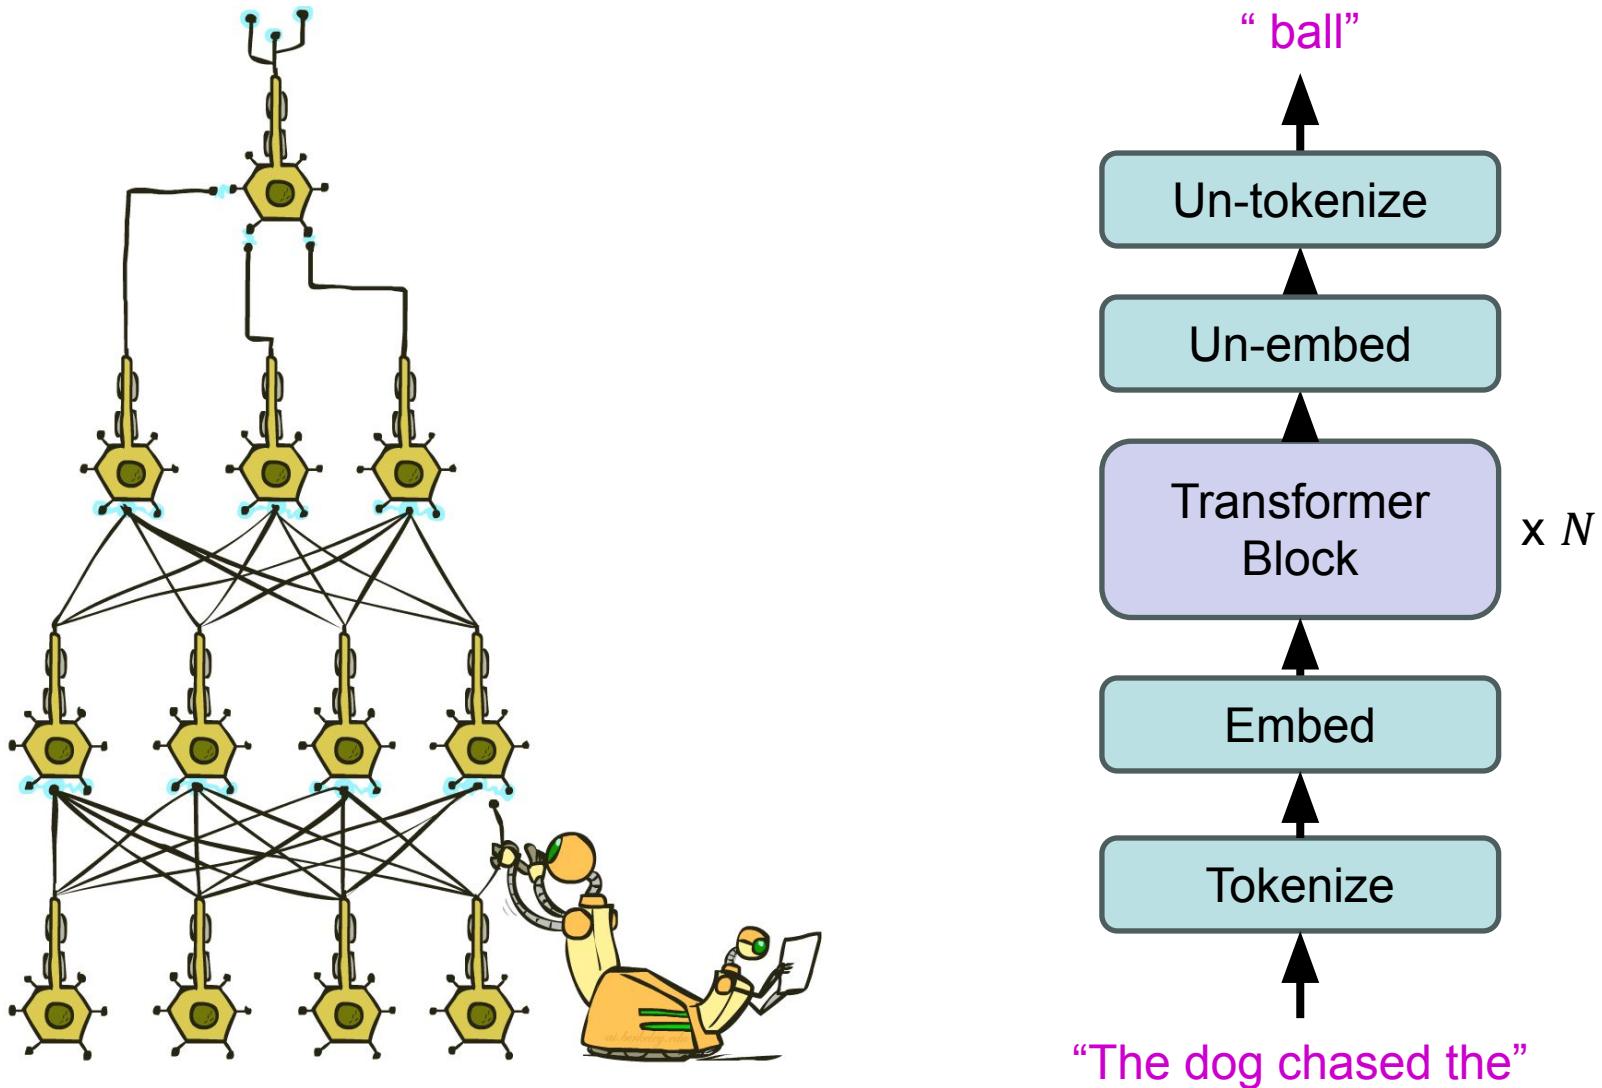
---



# Transformer Architecture



# Transformer Architecture



# Large Language Models

---

- ~~Feature engineering~~
  - Text tokenization
  - Word embeddings
- ~~Deep neural networks~~
  - Autoregressive models
  - Self attention mechanisms
  - Transformer architectures
- ~~Multi-class classification~~
  - Supervised learning
    - Self-supervised learning
    - Instruction tuning
  - Reinforcement learning
    - ... from human feedback (RLHF)

# Unsupervised / Self-Supervised Learning

---

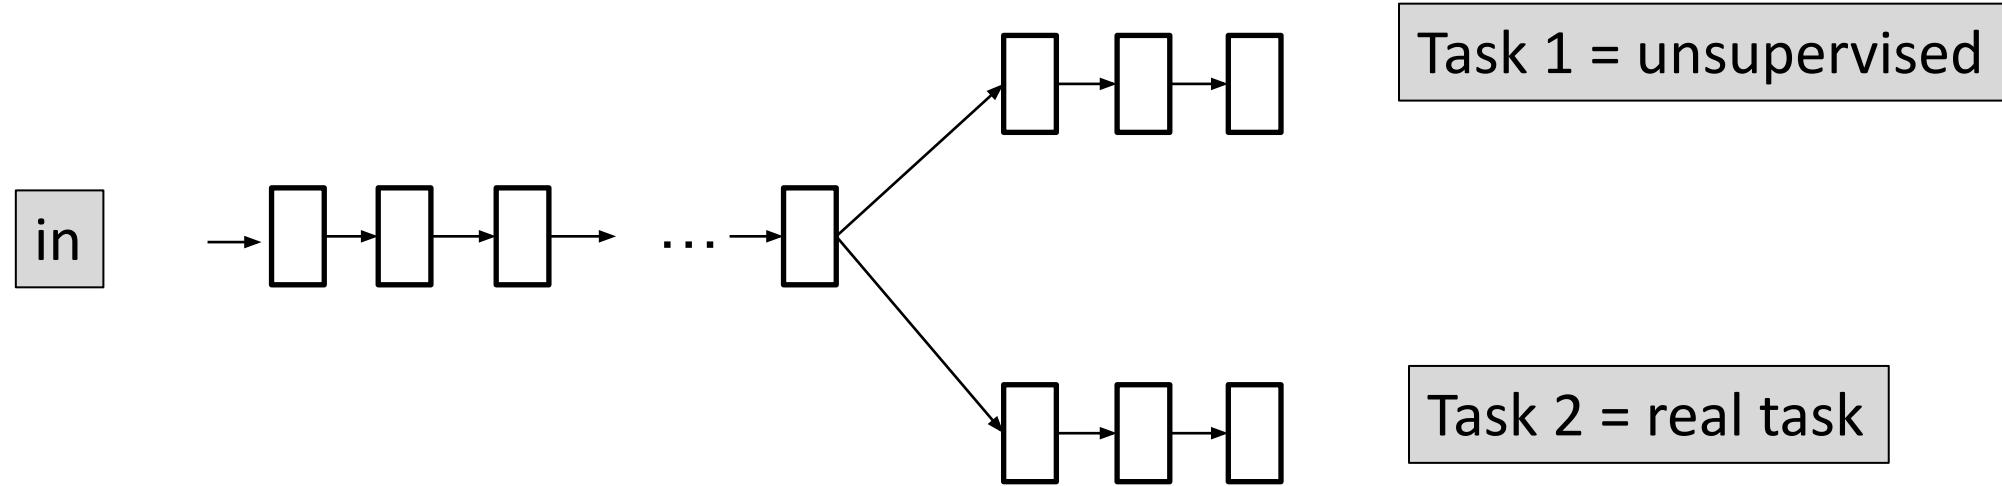
- Do we always need human supervision to learn features?
- Can't we learn general-purpose features?
- Key hypothesis:

Task 1 IF neural network smart enough to predict:

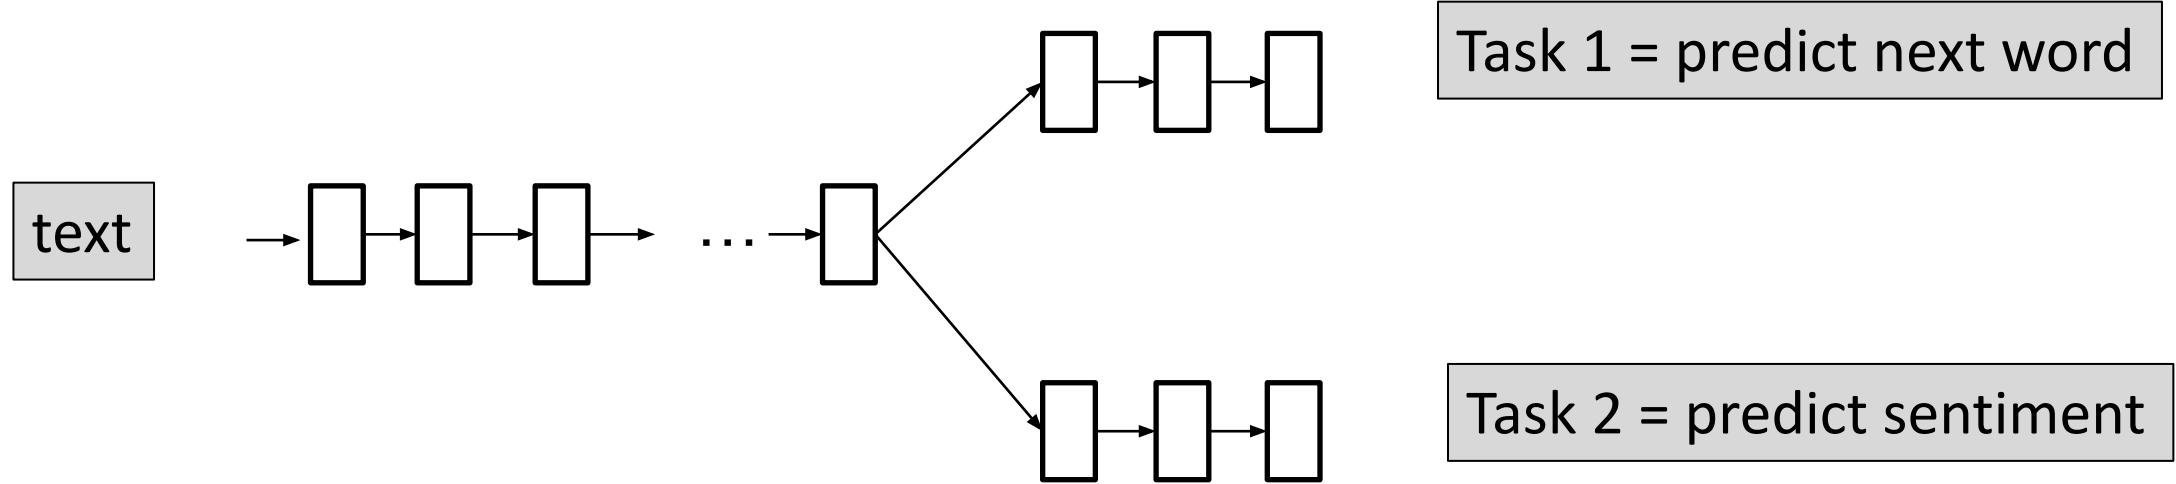
- Next frame in video
- Next word in sentence
- Generate realistic images
- ``Translate'' images
- ...

Task 2 THEN same neural network is ready to do Supervised Learning from a very small data-set

# Transfer from Unsupervised Learning



# Example Setting



# Image Pre-Training: Predict Missing Patch

---



# Pre-Training and Fine-Tuning

---

1

**Pre-Train:** train a large model with a lot of data on a self-supervised task

- Predict next word / patch of image
- Predict missing word / patch of image
- Predict if two images are related (contrastive learning)

2

**Fine-Tune:** continue training the same model on task you care about

# Instruction Tuning

---

- Task 1 = predict next word **(learns to mimic human-written text)**
  - Query: “What is population of Berkeley?”
  - Human-like completion: “This question always fascinated me!”
- Task 2 = generate **helpful** text
  - Query: “What is population of Berkeley?”
  - Helpful completion: “It is 117,145 as of 2021 census.”
- Fine-tune on collected examples of helpful human conversations
- Also can use Reinforcement Learning

# Reinforcement Learning from Human Feedback

---

## ■ MDP:

- **State:** sequence of words seen so far (ex. "What is population of Berkeley? ")
  - $100,000^{1,000}$  possible states
  - Huge, but can be processed with feature vectors or neural networks
- **Action:** next word (ex. "It", "chair", "purple", ...) (so 100,000 actions)
  - Hard to compute  $\max_a Q(s', a)$  when max is over 100K actions!
- **Transition T:** easy, just append action word to state words
  - s: "My name" a: "is" s': "My name is"
- **Reward R:** ???
  - Humans rate model completions (ex. "What is population of Berkeley? ")
    - "It is 117,145": +1
    - "It is 5": -1
    - "Destroy all humans": -1
  - Learn a reward model  $\hat{R}$  and use that (model-based RL)

# Knowing what to optimize for is very hard

- Clearly, we don't just want to predict the next word in internet text
- But even human feedback can have surprising / bad consequences
  - Sycophancy
  - Overconfidence
  - Length
  - ...



# Knowing what to optimize for is very hard

---

- More generally: lots of bad things happen when there is a gap between *what we really want to optimize for* and *what we train the model to optimize for*

Desired target	Actual target	Bias
Patient health needs	Patient health costs	Disparities in access to care
Severity of knee osteoarthritis	Severity as assessed by radiologist	Radiologists overlook features affecting underserved populations
Crime rates	Arrest rates	Disparities in policing
What users value	What users click on	Clickbait

# Large Language Models

---

- ~~Feature engineering~~

- ~~Text tokenization~~
- ~~Word embeddings~~

- ~~Deep neural networks~~

- ~~Autoregressive models~~
- ~~Self attention mechanisms~~
- ~~Transformer architectures~~

- ~~Multi-class classification~~

- ~~Supervised learning~~

- ~~Self-supervised learning~~
- ~~Instruction tuning~~

- ~~Reinforcement learning~~

- ~~... from human feedback (RLHF)~~

# Language models build a structured concept space

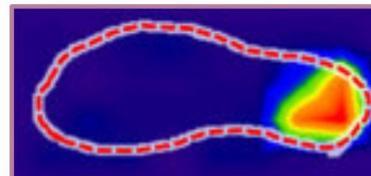
---



# Can other data (images/audio/...) be put in this space?



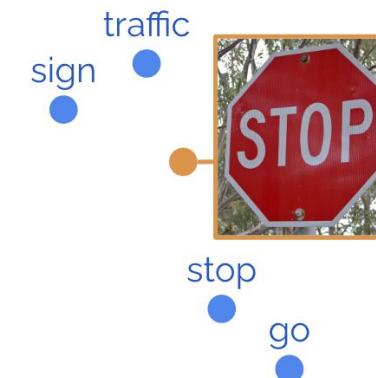
river  
water  
ocean



heel  
head  
toe



upward  
airplane



# Can we build a single model of all data types?

If



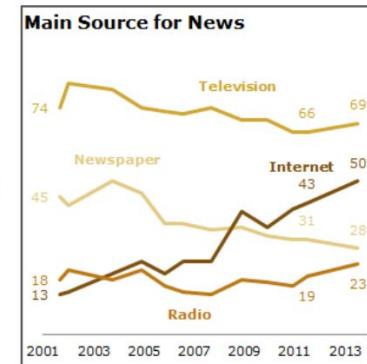
was invented by Wright brothers. Who invented

example from [Tsimpoukelli et al, 2021]



?

What is the fastest-growing news source according to



?

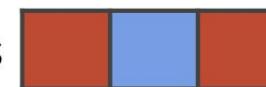
If



changes into

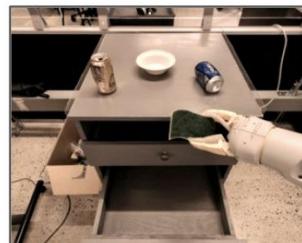


what does



change into?

What action should I take from



to accomplish “”?



# Can we build a single model of all data types?

## Mobile Manipulation



Human: Bring me the rice chips from the drawer. Robot: 1. Go to the drawers, 2. Open top drawer. I see . 3. Pick the green rice chip bag from the drawer and place it on the counter.

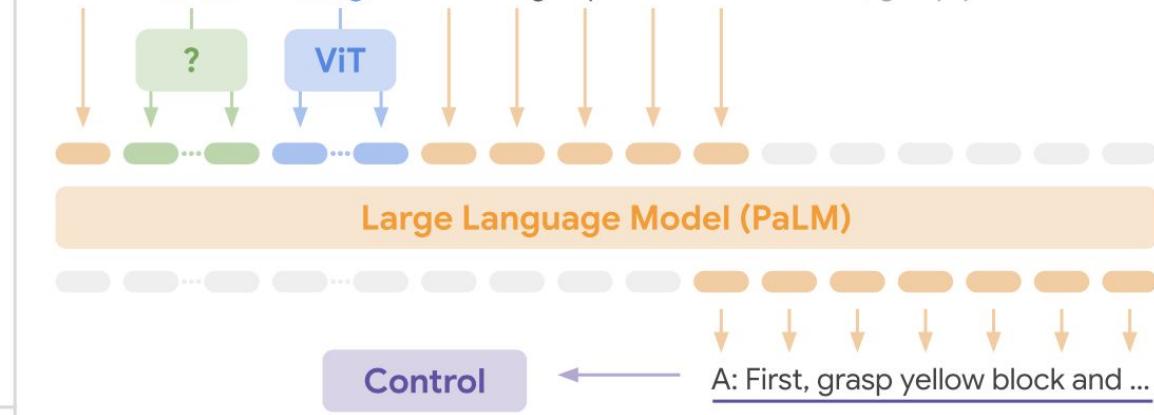
## Visual Q&A, Captioning ...



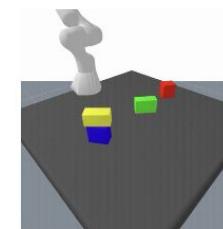
Given . Q: What's in the image? Answer in emojis.  
A:

## PaLM-E: An Embodied Multimodal Language Model

Given **<emb>** ... **<img>** Q: How to grasp blue block? A: First, grasp yellow block

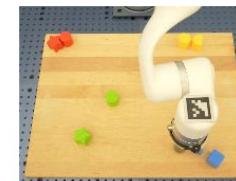


## Task and Motion Planning



Given **<emb>** Q: How to grasp blue block?  
A: First grasp yellow block and place it on the table, then grasp the blue block.

## Tabletop Manipulation



Given Task: Sort colors into corners.  
Step 1. Push the green star to the bottom left.  
Step 2. Push the green circle to the green star.

## Language Only Tasks



Describe the following   
A dog jumping over a hurdle at a dog show.

Q: Miami Beach borders which ocean? A: Atlantic. Q: What is  $372 \times 18$ ? A: 6696. Q: Write a Haiku about embodied LLMs. A: Embodied language. Models learn to understand.  
The world around them.